

# DOPP\_G18\_Ex3\_20210126\_Training200\_All\_Decades

January 26, 2021

## Table of Contents

### 1 Introduction to Assignment

#### 1.1 First draft

##### 1.1.1 Topic and Questions to answer

##### 1.1.2 Justification For Limit Of Scope

##### 1.1.3 Workflow plan & Project management

##### 1.1.4 Data

#### 1.2 Second Draft

#### 1.3 Pivoting Point

#### 1.4 Language change in Icelandic Parliamentary Speeches

### 2 Estimating publication year from Project Gutenberg

#### 2.1 Setup

##### 2.1.1 Import packages

##### 2.1.2 Define Constants

#### 2.2 Importing the data

##### 2.2.1 Getting the content

##### 2.2.2 Data Cleansing

###### 2.2.2.1 Read a single file

###### 2.2.2.2 Return list of all words

#### 2.3 Statistics

##### 2.3.1 First attempt

##### 2.3.2 Read all files, and do preprocessing

##### 2.3.3 Compare Word ranking between titles

#### 2.4 Second testing

##### 2.4.1 Read in from the decades files, and see the distributions

##### 2.4.2 Preliminary Conclusion

- 2.4.3 Compare ranking between upload-decades
- 2.5 Trying to fit models to predict
  - 2.5.1 Read in files
  - 2.5.2 Train models
- 2.6 Realisation and conclusion
- 3 Studying language change in Icelandic parliamentary speeches
  - 3.1 Introduction
  - 3.2 The task
  - 3.3 Setup
    - 3.3.1 Load required libraries
    - 3.3.2 Get the data
    - 3.3.3 Preprocessing helpers
  - 3.4 Preliminary Data Analysis
    - 3.4.1 Zipf's Law
    - 3.4.2 Disappearing words / new words
    - 3.4.3 Development of average sentence length
    - 3.4.4 n-grams
  - 3.5 Building model for classifying speeches
    - 3.5.1 Constructing training and test data
      - 3.5.1.1 Train data
      - 3.5.1.2 Test data
      - 3.5.1.3 See classes distribution within train and test sets
    - 3.5.2 Text feature extraction
      - 3.5.2.1 TF-IDF
      - 3.5.2.2 Word2Vec
      - 3.5.2.3 Doc2Vec
    - 3.5.3 Classifiers
      - 3.5.3.1 Multinomial Naive Bayes (MNB)
      - 3.5.3.2 Support Vector Machines
      - 3.5.3.3 Random Forest Classifier
    - 3.5.4 Train models
      - 3.5.4.1 Model 1: TFIDF vectorizer, select K best and Multinomial Naive Bayes

- 3.5.4.2 Model 2: TFIDF vectorizer, select K best and SVC
- 3.5.4.3 Model 3: TFIDF vectorizer, select K best and Random Forest Classifier
- 3.5.4.4 Model 4: Word2Vec and SVC
- 3.5.4.5 Model 5: Word2Vec and Random Forest Classifier
- 3.5.4.6 Model 6: Doc2Vec and Support Vector Machines
- 3.5.4.7 Model 7: Doc2Vec and Random Forest Classifier
- 3.5.5 Compare CV results from trained models
  - 3.5.5.1 Raw results
  - 3.5.5.2 Tradeoff score vs mean fit time
  - 3.5.5.3 Best estimator from each model
  - 3.5.5.4 Next steps
- 3.6 Evaluation and model selection
- 4 Conclusion
- 5 Further Works
  - 5.1 Predict Different Sources
  - 5.2 Treating years as Continuous Variables
  - 5.3 Gaining insight into Explanatory Variables
  - 5.4 Additional Feature Extraction and Classifiers

## 1 Introduction to Assignment

This is the third Exercise of **188.995 Data-Oriented Programming Paradigms**

We are group 18, and consist of: \* Guillermo Alamán Requena, Matr. Nr: 11937906 \* Michael Ferdinand Moser, Matr. Nr: 01123077 \* Paul Joe Maliakel, Matr. Nr: 12012422 \* Gunnar Sjúrdarson Knudsen, Matr. Nr: 12028205

In this task we were asked to choose one vaguely worded question, and then narrow the scope, figuring out how to get the data, before finally solving the question at hand. We chose **Question 21**, which contains: \* How does the use of various communication languages in countries change over time? \* Which languages grow and which disappear, and what are their characteristics? \* Are there other factors that correlate with the appearance or disappearance of languages?

We soon realized that the question as stated is far too broad, and we therefore had to limit it.

After having discussed among our groups, we came to the following plan:

## 1.1 First draft

### 1.1.1 Topic and Questions to answer

We've selected question 21, which is regarding how communication languages in countries change over time.

After having discussed the data available, and planned a workflow, we've decided to try to answer the questions: \* How has the English language changed in the past 100 years based on word frequencies, sentence length, ...? \* Can we find parallel developments between different genres of text? \* Can the publication year of a movie/article/whatever be predicted based on the text and its characteristics?

### 1.1.2 Justification For Limit Of Scope

The sample questions stated in the task description are too broad, to be answered in a single 160 hour project. \* Lot's of issues, such as: \* Lack of census data; \* other changes such as phonetic, semantic and syntactic meanings; \* High correlation with e.g: \* country population \* age of speakers \* ... \* What counts as a language? \* dialect? \* Mutually Intelligible? \* Political dimensions \* Multilingual people \* How do we check accuracy of the available data? \* ...

Historical data for language use is likely not available for most languages, as it's topics for great research to estimate merely historical populations - especially before 1850 or so. The evolution of languages are much less documented. Lack of census data overall, but other changes are even harder to gauge, such as phonetic, semantic, and syntactic meanings. Highly correlated with population of countries, but also with "hidden" correlations, such as age of speakers, ... Even dead languages can be revived.

What constitutes a language? Dialect? Mutually Intelligible? Also do not forget the political dimension, e.g. Croatian/Serbian really are just dialects of the same language but they want to keep separate. On the other end of this scheme the variant of Chinese spoken in Beijing may be drastically different from the Chinese spoken in other regions of the country, but still falls under the same "Chinese" umbrella to communicate unity.

How much is spoken? Should we consider people who studied a language as their second, third... language? If so, how well should be the command over the language for the person to count? A1/B1/C2 level? %How do we check accuracy of the available data?

### 1.1.3 Workflow plan & Project management

- Outline the plan
  - Get, understand and clean data: articles/movie scripts/video transcripts over the years (see next section)
  - Train-test split: keeping proportion of publication years within the splits.
  - Preprocessing: text feature extraction, feature selection, scaling, etc. (Come back here if necessary)
  - Visualization: evolution of words over the years, word-clouds and other relevant characteristics.
  - Define evaluation metrics, train different models/parameters using CV and select best one for predictions.
  - Predict, conclude, report and publish notebook in Kaggle Kernel.

- How the work will be divided up between group members
  - Acquisition, cleaning and preprocessing of the data will be done commonly.
  - Each member of the group will train a model and report results using same evaluation metrics.
  - Jointly choose the best model and conclude.
  - Presentation, report and publishing will be also split.
- Timeline: To be defined after review meeting

#### 1.1.4 Data

Our goal is to get a dataset similar to:

Corpus	Year Published	Type	...
Text1	1976	News	...
Text2	1976	Movie Script	...
...	...	...	...
TextN	2009	Scientific Article	...

Feature extraction from texts will be performed to obtain appropriate features for modeling. To build a dataset like this one, we will rely on the following kind sources:

- <https://www.kaggle.com/asad1m9a9h6mood/news-articles> - News articles from 2015 until date.
- <https://www.kaggle.com/snapcrack/all-the-news> - 143000 articles from 15 American Publications.
- NLTK
- ...

## 1.2 Second Draft

After having a preliminary meeting with Univ.Prof. Dr. Hanbury and Dipl.-Ing. Dr. Piroi, who gave great input, we decided to further limit our goal to only use Project Gutenberg as a data source, and setting our hypothesis to see whether it was possible to generate a model that predicted the publication year/decade for a set of books.

## 1.3 Pivoting Point

After having done a decent portion of work, we reached to the conclusion that our dataset was not suitable to solve the question we had originally set out, and we were forced to pivot.

We discussed whether we wanted to change the goal from classifying, but as we were all quite interested in a classification algorithm, and wanted to do proper NLP, we instead searched for another dataset.

## 1.4 Language change in Icelandic Parliamentary Speeches

We found the dataset with all Icelandic parliamentary speeches going back a century. This is further described in section 3. With this great dataset, our goal was to develop a model that could try to predict which decade a speech is from.

## 2 Estimating publication year from Project Gutenberg

This was the attempt at our first hypothesis. We import a large corpus of books from Project Gutenberg, and cleanse the data, so it's ready for machine learning

### 2.1 Setup

We start by setting up all packages needed for the project

#### 2.1.1 Import packages

```
[1]: from __future__ import absolute_import
from builtins import str
import os
from six import u

from os import listdir
from os.path import isfile, join

import nltk
import re
from operator import itemgetter
import pandas as pd
from functools import reduce

import random

pd.set_option('display.max_rows', None)

import math

from sklearn.feature_extraction.text import TfidfTransformer

from pprint import pprint
from time import time
import logging

from sklearn.datasets import fetch_20newsgroups
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.linear_model import SGDClassifier
from sklearn.naive_bayes import MultinomialNB
from sklearn.model_selection import GridSearchCV
from sklearn.pipeline import Pipeline

from sklearn.feature_selection import SelectKBest, chi2
```

### 2.1.2 Define Constants

Constant that are used in this part is also set

```
[2]: file_path = "processedData"

TEXT_START_MARKERS = frozenset((u(_) for _ in (
    "*END*THE SMALL PRINT",
    "*** START OF THE PROJECT GUTENBERG",
    "*** START OF THIS PROJECT GUTENBERG",
    "This etext was prepared by",
    "E-text prepared by",
    "Produced by",
    "Distributed Proofreading Team",
    "Proofreading Team at http://www.pgdp.net",
    "http://gallica.bnf.fr)",
    "    http://archive.org/details/",
    "http://www.pgdp.net",
    "by The Internet Archive)",
    "by The Internet Archive/Canadian Libraries",
    "by The Internet Archive/American Libraries",
    "public domain material from the Internet Archive",
    "Internet Archive)",
    "Internet Archive/Canadian Libraries",
    "Internet Archive/American Libraries",
    "material from the Google Print project",
    "*END THE SMALL PRINT",
    "***START OF THE PROJECT GUTENBERG",
    "This etext was produced by",
    "*** START OF THE COPYRIGHTED",
    "The Project Gutenberg",
    "http://gutenberg.spiegel.de/ erreichbar.",
    "Project Runeberg publishes",
    "Beginning of this Project Gutenberg",
    "Project Gutenberg Online Distributed",
    "Gutenberg Online Distributed",
    "the Project Gutenberg Online Distributed",
    "Project Gutenberg TEI",
    "This eBook was prepared by",
    "http://gutenberg2000.de erreichbar.",
    "This Etext was prepared by",
    "This Project Gutenberg Etext was prepared by",
    "Gutenberg Distributed Proofreaders",
    "Project Gutenberg Distributed Proofreaders",
    "the Project Gutenberg Online Distributed Proofreading Team",
    "**The Project Gutenberg",
    "*SMALL PRINT!",
    "More information about this book is at the top of this file.",
```

```

    "tells you about restrictions in how the file may be used.",
    "l'authorization à les utiliser pour preparer ce texte.",
    "of the etext through OCR.",
    "*****These eBooks Were Prepared By Thousands of Volunteers!*****",
    "We need your donations more than ever!",
    " *** START OF THIS PROJECT GUTENBERG",
    "****      SMALL PRINT!",
    '["Small Print" V.',
    '      (http://www.ibiblio.org/gutenberg/',
    'and the Project Gutenberg Online Distributed Proofreading Team',
    'Mary Meehan, and the Project Gutenberg Online Distributed Proofreading',
    '      this Project Gutenberg edition.',
    )))

```

```

TEXT_END_MARKERS = frozenset((u(_) for _ in (
    "*** END OF THE PROJECT GUTENBERG",
    "*** END OF THIS PROJECT GUTENBERG",
    "***END OF THE PROJECT GUTENBERG",
    "End of the Project Gutenberg",
    "End of The Project Gutenberg",
    "Ende dieses Project Gutenberg",
    "by Project Gutenberg",
    "End of Project Gutenberg",
    "End of this Project Gutenberg",
    "Ende dieses Projekt Gutenberg",
    "      ***END OF THE PROJECT GUTENBERG",
    "*** END OF THE COPYRIGHTED",
    "End of this is COPYRIGHTED",
    "Ende dieses Etextes ",
    "Ende dieses Project Gutenber",
    "Ende diese Project Gutenberg",
    "**This is a COPYRIGHTED Project Gutenberg Etext, Details Above**",
    "Fin de Project Gutenberg",
    "The Project Gutenberg Etext of ",
    "Ce document fut presente en lecture",
    "Ce document fut présenté en lecture",
    "More information about this book is at the top of this file.",
    "We need your donations more than ever!",
    "END OF PROJECT GUTENBERG",
    " End of the Project Gutenberg",
    " *** END OF THIS PROJECT GUTENBERG",
    )))

```

```

LEGALESE_START_MARKERS = frozenset((u(_) for _ in (
    "<<THIS ELECTRONIC VERSION OF",

```



```

)))

LEGALESE_END_MARKERS = frozenset((u(_) for _ in (
    "SERVICE THAT CHARGES FOR DOWNLOAD",
)))

TITLE_MARKERS = frozenset((u(_) for _ in (
    "Title:",
)))

AUTHOR_MARKERS = frozenset((u(_) for _ in (
    "Author:",
)))

DATE_MARKERS = frozenset((u(_) for _ in (
    "Release Date:", "Release Date:"
)))

LANGUAGE_MARKERS = frozenset((u(_) for _ in (
    "Language:",
)))

ENCODING_MARKERS = frozenset((u(_) for _ in (
    "Character set encoding:",
)))

```

## 2.2 Importing the data

This is a very rough first draft at importing and cleansing the data. Solution is heavily inspired by <https://gist.github.com/mbforbes/cee3fd5bb3a797b059524fe8c8ccdc2b>

### 2.2.1 Getting the content

Start by downloading the repository of (english) books. This is done in bash. Only tested on Ubuntu, but mac should work the same

```
wget -m -H -nd "http://www.gutenberg.org/robot/harvest?filetypes[]=txt&langs[]=en"
```

```
http://www.gutenberg.org/robot/harvest?offset=40532&filetypes[]=txt&langs[]=en
```

Takes a few hours to run, and is stored in a folder called rawContent. This is then copied to another folder, and we can start to clean up the mess

First we delete some duplications of the same books:

```
ls | grep "\-8.zip" | xargs rm
ls | grep "\-0.zip" | xargs rm
```

We can then unzip the files, and remove the zip files

```
unzip "*.zip"
rm *.zip
```

Next we take care of some nested foldering

```
mv */*.txt ./
```

And finally, we remove all rubbish that isn't a real book:

```
ls | grep -v "\.txt" | xargs rm -rf
```

## 2.2.2 Data Cleansing

As the data is not given in a computer-friendly format, a lot of string operations are needed

### Read a single file

```
[3]: def read_file(file_name):
    file = open(file_name, encoding="ISO-8859-1")
    file_content = file.read()

    lines = file_content.splitlines()
    sep = str(os.linesep)

    # Initialize results for single book
    content_lines = []
    i = 0
    footer_found = False
    ignore_section = False

    title = ""
    author = ""
    date = ""
    language = ""
    encoding = ""
    year = 0

    # Reset flags for each book
    title_found = False
    author_found = False
    date_found = False
    language_found = False
    encoding_found = False

    for line in lines:
        reset = False

        #print(line)
        if i <= 600:
            # Shamelessly stolen
            if any(line.startswith(token) for token in TEXT_START_MARKERS):
                reset = True
```

```

# Extract Metadata
if title_found == False:
    if any(line.startswith(token) for token in TITLE_MARKERS):
        title_found = True
        title = line
if author_found == False:
    if any(line.startswith(token) for token in AUTHOR_MARKERS):
        author_found = True
        author = line
if date_found == False:
    if any(line.startswith(token) for token in DATE_MARKERS):
        date_found = True
        date = line
        year = int(re.findall(r'\d{4}', date)[0])
if language_found == False:
    if any(line.startswith(token) for token in_
↪LANGUAGE_MARKERS):
        language_found = True
        language = line
if encoding_found == False:
    if any(line.startswith(token) for token in_
↪ENCODING_MARKERS):
        encoding_found = True
        encoding = line

# More theft from above
if reset:
    content_lines = []
    continue

# I feel like a criminal by now. Guess what? Also stolen
if i >= 100:
    if any(line.startswith(token) for token in TEXT_END_MARKERS):
        footer_found = True

    if footer_found:
        break

if any(line.startswith(token) for token in LEGALESE_START_MARKERS):
    ignore_section = True
    continue
elif any(line.startswith(token) for token in LEGALESE_END_MARKERS):
    ignore_section = False
    continue

if not ignore_section:
    if line != "": # Screw the blank lines

```

```

        content_lines.append(line.rstrip(sep))
        i += 1

    sep.join(content_lines)

# Do more cleaning
for token in TITLE_MARKERS:
    title = title.replace(token, '').lstrip().rstrip()
for token in AUTHOR_MARKERS:
    author = author.replace(token, '').lstrip().rstrip()
for token in LANGUAGE_MARKERS:
    language = language.replace(token, '').lstrip().rstrip()
for token in DATE_MARKERS:
    date = date.replace(token, '').lstrip().rstrip()
for token in ENCODING_MARKERS:
    encoding = encoding.replace(token, '').lstrip().rstrip()
return title, author, date, year, language, encoding, content_lines

```

**Return list of all words** Currently quite an empty function. However, I assume that some cleaning of cases etc. will be done here

```

[4]: def get_words(content_lines):
    all_text_lower = " ".join(content_lines).lower()
    words = re.findall(r'(\b[A-Za-z][a-z]{2,9}\b)', all_text_lower)

    # Do more cleansing. E.g. cases and stuff

    return words

```

## 2.3 Statistics

We start by doing some exploratory data analysis, to see how well our scraping works

### 2.3.1 First attempt

Trying a simple word frequency

```

[5]: def get_word_frequencies(words):
    frequency = {}
    for word in words:
        count = frequency.get(word, 0)
        frequency[word] = count + 1

    word_count = len(words)
    unique_word_count = 0
    word_list = []
    word_list_count = []

```

```

for key, value in reversed(sorted(frequency.items(), key = itemgetter(1))):
    word_list.append(key)
    word_list_count.append(value)
    unique_word_count = unique_word_count + 1

word_list_freq = [freq / word_count for freq in word_list_count]

word_freq = pd.DataFrame(list(zip(word_list, word_list_count,
↪word_list_freq))
                        , columns = ['Word', 'count', 'freq'])

word_freq['rank'] = word_freq['count'].rank(ascending = False,
↪method="dense")

return(word_freq, unique_word_count)

```

### 2.3.2 Read all files, and do preprocessing

Well... Only ten files currently

```

[6]: # Get all filenames
files = [f for f in listdir(file_path) if isfile(join(file_path, f))]
files = list(filter(lambda file: file[0].isdigit(), files))
random.shuffle(files)

# Do only subset
files = files[0:10]

list_of_file = []
list_of_title = []
list_of_author = []
list_of_date = []
list_of_year = []
list_of_language = []
list_of_encoding = []
list_of_word_count = []
list_of_unique_word_count = []
list_of_word_frequencies = []
iter_ = 0

for file in files:
    # Read in basic information from file
    title, author, date, year, language, encoding, content_lines =
↪read_file(file_path + "/" + file)
    line_count = len(content_lines)

```

```

# Not sure if we want this for later:
#content_all = " ".join(content_lines)

# Split into words (and do various cleaning)
words = get_words(content_lines)
word_count = len(words)

# First analysis, but should do something proper
word_frequencies_table, unique_word_count = get_word_frequencies(words)

# Append to results
list_of_file.append(file)
list_of_title.append(title)
list_of_author.append(author)
list_of_date.append(date)
list_of_year.append(year)
list_of_language.append(language)
list_of_encoding.append(encoding)
list_of_word_count.append(word_count)
list_of_unique_word_count.append(unique_word_count)
list_of_word_frequencies.append(word_frequencies_table)

# Show basic information
#print(iter_)
iter_ = iter_ + 1
#print("#####")
#print("#####")
#print("Filename: " + str(file))
#print("Title: " + str(title))
#print("Author(s): " + str(author))
#print("Date: " + str(date))
#print("Year: " + str(year))
#print("Language: " + str(language))
#print("Encoding: " + str(encoding))
#print("#####")
#print("Words in book: " + str(word_count))
#print("Unique words in book: " + str(unique_word_count))
#print("#####")
#print(word_frequencies_table)

# Feel free to change to dict? list? separate files?
## nested dataframes works, but looks super ugly when printing
### Fuck it - This is tooo useless killing it again
#all_res = pd.DataFrame(list(zip(list_of_file
#                               , list_of_title
#                               , list_of_author

```

```

#             , list_of_date
#             , list_of_language
#             , list_of_encoding
#             , list_of_word_count
#             , list_of_unique_word_count
#             , list_of_word_frequencies
#         ))
#         , columns = ['file'
#                     , 'title'
#                     , 'author'
#                     , 'date'
#                     , 'language'
#                     , 'encoding'
#                     , 'word_count'
#                     , 'unique_word_count'
#                     , 'word_frequencies'
#                 ]
#     )
#

```

### 2.3.3 Compare Word ranking between titles

This is our first attempt at seeing how the ranking of words change between titles. Idea is to see that the zipf-distribution changes as time passes buy

```

[7]: list_count= []
list_freq = []
list_rank = []

col_names = list_of_title.copy()
col_names.insert(0,'Word')

for df in list_of_word_frequencies:
    list_count.append(df[['Word', 'count']])
    list_freq.append(df[['Word', 'freq']])
    list_rank.append(df[['Word', 'rank']])

df_count = reduce(lambda left, right: pd.merge(left, right, on="Word",
    ↪how='outer'), list_count)
df_count.columns = col_names
df_count['Sum'] = df_count.drop('Word', axis=1).apply(lambda x: x.sum(), axis=1)
df_count = df_count.sort_values(ascending = False, by=['Sum'])

df_freq = reduce(lambda left, right: pd.merge(left, right, on="Word",
    ↪how='outer'), list_freq)
df_freq.columns = col_names

```

```
df_freq['Avg'] = df_freq.drop('Word', axis=1).apply(lambda x: x.mean(), axis=1)
df_freq = df_freq.sort_values(ascending = False, by=['Avg'])

df_rank = reduce(lambda left, right: pd.merge(left, right, on="Word",
→how='outer'), list_rank)
df_rank.columns = col_names
df_rank['Avg'] = df_rank.drop('Word', axis=1).apply(lambda x: x.mean(), axis=1)
df_rank = df_rank.sort_values(by=['Avg'])
```

```
[8]: df_rank.head(30)
```

```
[8]:
```

	Word	Tonio, Son of the Sierras	The Modern Railroad	\
0	the	1.0	1.0	1.0
1	and	2.0	2.0	2.0
17417	gutenberg	NaN	NaN	NaN
17418	donations	NaN	NaN	NaN
19892	bors	NaN	NaN	NaN
2	that	3.0	6.0	3.0
5	was	6.0	3.0	6.0
4	for	5.0	8.0	4.0
3	with	4.0	7.0	10.0
17419	org	NaN	NaN	NaN
6	his	7.0	5.0	7.0
17420	ebooks	NaN	NaN	NaN
17421	donate	NaN	NaN	NaN
17424	email	NaN	NaN	NaN
6510	willet	NaN	11.0	NaN
17427	exempt	NaN	NaN	NaN
17426	pglaf	NaN	NaN	NaN
8	but	9.0	10.0	20.0
17450	fairbanks	NaN	NaN	NaN
17453	ein	NaN	NaN	NaN
17428	newsletter	NaN	NaN	NaN
17429	editions	NaN	NaN	NaN
17431	concept	NaN	NaN	NaN
17432	originator	NaN	NaN	NaN
17435	addresses	NaN	NaN	NaN
17436	donation	NaN	NaN	NaN
17437	swamp	NaN	NaN	NaN
17439	donors	NaN	NaN	NaN
17452	deductible	NaN	NaN	NaN
17440	fees	NaN	NaN	NaN

	The Home Mission	The Everlasting Mercy	\
0	1.0	1.0	
1	2.0	2.0	
17417	NaN	NaN	



17418	NaN	NaN
19892	NaN	NaN
2	5.0	6.0
5	4.0	8.0
4	12.0	9.0
3	10.0	5.0
17419	NaN	NaN
6	8.0	7.0
17420	NaN	NaN
17421	NaN	NaN
17424	NaN	NaN
6510	NaN	NaN
17427	NaN	NaN
17426	NaN	NaN
8	13.0	12.0
17450	NaN	NaN
17453	NaN	NaN
17428	NaN	NaN
17429	NaN	NaN
17431	NaN	NaN
17432	NaN	NaN
17435	NaN	NaN
17436	NaN	NaN
17437	NaN	NaN
17439	NaN	NaN
17452	NaN	NaN
17440	NaN	NaN

Chambers's Journal of Popular Literature, Science, and Art, No. 725,  
November 17, 1877 \

0	1.0
1	2.0
17417	NaN
17418	NaN
19892	NaN
2	3.0
5	4.0
4	5.0
3	10.0
17419	NaN
6	8.0
17420	NaN
17421	NaN
17424	NaN
6510	NaN
17427	NaN
17426	NaN

8	11.0
17450	NaN
17453	NaN
17428	NaN
17429	NaN
17431	NaN
17432	NaN
17435	NaN
17436	NaN
17437	NaN
17439	NaN
17452	NaN
17440	NaN

	Eternal Father, Strong to Save	The Masked Bridal	The Farmer's Boy \
0	1.0	1.0	1.0
1	2.0	2.0	2.0
17417	3.0	NaN	NaN
17418	4.0	NaN	NaN
19892	NaN	NaN	NaN
2	12.0	6.0	5.0
5	NaN	7.0	13.0
4	10.0	10.0	7.0
3	7.0	8.0	4.0
17419	8.0	NaN	NaN
6	NaN	11.0	3.0
17420	9.0	NaN	NaN
17421	10.0	NaN	NaN
17424	11.0	NaN	NaN
6510	NaN	NaN	NaN
17427	12.0	NaN	NaN
17426	12.0	NaN	NaN
8	12.0	14.0	12.0
17450	13.0	NaN	NaN
17453	13.0	NaN	NaN
17428	13.0	NaN	NaN
17429	13.0	NaN	NaN
17431	13.0	NaN	NaN
17432	13.0	NaN	NaN
17435	13.0	NaN	NaN
17436	13.0	NaN	NaN
17437	13.0	NaN	NaN
17439	13.0	NaN	NaN
17452	13.0	NaN	NaN
17440	13.0	NaN	NaN

Talents, Incorporated	Avg
-----------------------	-----

0	1.0	1.000000
1	2.0	2.000000
17417	NaN	3.000000
17418	NaN	4.000000
19892	5.0	5.000000
2	4.0	5.300000
5	3.0	6.000000
4	7.0	7.700000
3	15.0	8.000000
17419	NaN	8.000000
6	19.0	8.333333
17420	NaN	9.000000
17421	NaN	10.000000
17424	NaN	11.000000
6510	NaN	11.000000
17427	NaN	12.000000
17426	NaN	12.000000
8	8.0	12.100000
17450	NaN	13.000000
17453	NaN	13.000000
17428	NaN	13.000000
17429	NaN	13.000000
17431	NaN	13.000000
17432	NaN	13.000000
17435	NaN	13.000000
17436	NaN	13.000000
17437	NaN	13.000000
17439	NaN	13.000000
17452	NaN	13.000000
17440	NaN	13.000000

```
[9]: #df_freq['Avg'] = df_freq.drop('Word', axis=1).apply(lambda x: x.mean(), axis=1)
df_freq = df_freq.sort_values(ascending = False, by=['Avg'])

df_freq.head(20)
```

```
[9]:
```

	Word	Tonio, Son of the Sierras	The Modern Railroad \
0	the	0.083286	0.109511
1	and	0.052821	0.035380
17417	gutenberg	NaN	NaN
17418	donations	NaN	NaN
19892	bors	NaN	NaN
17419	org	NaN	NaN
2	that	0.010917	0.018102
5	was	0.008057	0.008843
17420	ebooks	NaN	NaN
6	his	0.007784	0.007617

3	with	0.010321	0.010634	0.007177
4	for	0.009251	0.009499	0.012415
11	her	0.005645	0.005885	0.000913
7	you	0.007212	0.005899	0.003580
17421	donate	NaN	NaN	NaN
33	had	0.002611	0.017474	0.003669
13	not	0.004949	0.005677	0.004101
17424	email	NaN	NaN	NaN
8	but	0.006416	0.007269	0.004005
9	from	0.006317	0.005968	0.007345

	The Home Mission	The Everlasting Mercy \
0	0.059477	0.068943
1	0.039817	0.066598
17417	NaN	NaN
17418	NaN	NaN
19892	NaN	NaN
17419	NaN	NaN
2	0.015689	0.010037
5	0.016688	0.008536
17420	NaN	NaN
6	0.012575	0.008723
3	0.010935	0.010412
4	0.010103	0.007316
11	0.022654	0.003564
7	0.013502	0.013226
17421	NaN	NaN
33	0.010388	0.003658
13	0.011125	0.003002
17424	NaN	NaN
8	0.009009	0.006003
9	0.006513	0.006097

Chambers's Journal of Popular Literature, Science, and Art, No. 725,  
November 17, 1877 \

0	0.070279
1	0.041293
17417	NaN
17418	NaN
19892	NaN
17419	NaN
2	0.016969
5	0.012454
17420	NaN
6	0.009176
3	0.008230
4	0.011725

11	0.011652
7	0.010342
17421	NaN
33	0.008958
13	0.007647
17424	NaN
8	0.008084
9	0.005608

	Eternal Father, Strong to Save	The Masked Bridal	The Farmer's Boy \
0	0.056075	0.061594	0.082150
1	0.042056	0.037046	0.048715
17417	0.032710	NaN	NaN
17418	0.028037	NaN	NaN
19892	NaN	NaN	NaN
17419	0.014019	NaN	NaN
2	0.004673	0.019774	0.012615
5	NaN	0.016346	0.004923
17420	0.011682	NaN	NaN
6	NaN	0.010248	0.020255
3	0.016355	0.013447	0.013435
4	0.009346	0.011631	0.007282
11	NaN	0.028121	0.004769
7	NaN	0.020796	0.000462
17421	0.009346	NaN	NaN
33	NaN	0.012124	0.002923
13	0.014019	0.008419	0.004974
17424	0.007009	NaN	NaN
8	0.004673	0.008335	0.004974
9	0.009346	0.005641	0.009641

	Talents, Incorporated	Avg
0	0.074133	0.074600
1	0.029507	0.044609
17417	NaN	0.032710
17418	NaN	0.028037
19892	0.014087	0.014087
17419	NaN	0.014019
2	0.015294	0.013837
5	0.019294	0.012707
17420	NaN	0.011682
6	0.006389	0.011057
3	0.007547	0.010849
4	0.011295	0.009986
11	0.001811	0.009446
7	0.009458	0.009386
17421	NaN	0.009346

33	0.008729	0.007837
13	0.008175	0.007209
17424	NaN	0.007009
8	0.010414	0.006918
9	0.006113	0.006859

```
[10]: df_count.head(20)
```

```
[10]:
```

	Word	Tonio, Son of the Sierras	The Modern Railroad \
0	the	3349.0	5818.0
1	and	2124.0	3817.0
2	that	439.0	1033.0
5	was	324.0	1388.0
4	for	372.0	686.0
6	his	313.0	1209.0
3	with	415.0	768.0
11	her	227.0	425.0
7	you	290.0	426.0
33	had	105.0	1262.0
16	she	164.0	258.0
8	but	258.0	525.0
9	from	254.0	431.0
13	not	199.0	410.0
19	have	158.0	311.0
21	this	151.0	260.0
22	they	144.0	430.0
15	there	175.0	398.0
10	all	240.0	392.0
44	him	84.0	575.0

	The Home Mission	The Everlasting Mercy \
0	2502.0	735.0
1	1675.0	710.0
2	660.0	107.0
5	702.0	91.0
4	425.0	78.0
6	529.0	93.0
3	460.0	111.0
11	953.0	38.0
7	568.0	141.0
33	437.0	39.0
16	649.0	40.0
8	379.0	64.0
9	274.0	65.0
13	468.0	32.0
19	243.0	23.0
21	328.0	32.0

22	177.0	73.0
15	172.0	49.0
10	209.0	124.0
44	257.0	53.0

Chambers's Journal of Popular Literature, Science, and Art, No. 725,  
November 17, 1877 \

0	965.0
1	567.0
2	233.0
5	171.0
4	161.0
6	126.0
3	113.0
11	160.0
7	142.0
33	123.0
16	102.0
8	111.0
9	77.0
13	105.0
19	105.0
21	83.0
22	86.0
15	49.0
10	67.0
44	64.0

	Eternal Father, Strong to Save	The Masked Bridal	The Farmer's Boy \
0	24.0	5121.0	1602.0
1	18.0	3080.0	950.0
2	2.0	1644.0	246.0
5	NaN	1359.0	96.0
4	4.0	967.0	142.0
6	NaN	852.0	395.0
3	7.0	1118.0	262.0
11	NaN	2338.0	93.0
7	NaN	1729.0	9.0
33	NaN	1008.0	57.0
16	NaN	1903.0	32.0
8	2.0	693.0	97.0
9	4.0	469.0	188.0
13	6.0	700.0	97.0
19	2.0	772.0	98.0
21	1.0	434.0	130.0
22	NaN	193.0	71.0
15	NaN	260.0	38.0

10	2.0	312.0	76.0
44	NaN	517.0	82.0

	Talents, Incorporated	Sum
0	2947.0	36735.0
1	1173.0	18531.0
2	608.0	7232.0
5	767.0	6002.0
4	449.0	4834.0
6	254.0	4722.0
3	300.0	4450.0
11	72.0	4420.0
7	376.0	4128.0
33	347.0	3836.0
16	97.0	3287.0
8	414.0	3043.0
9	243.0	2922.0
13	325.0	2854.0
19	214.0	2511.0
21	119.0	2196.0
22	397.0	2168.0
15	394.0	2142.0
10	169.0	2048.0
44	122.0	2035.0

## 2.4 Second testing

This definately needs some proper refactoring, but Was curious whether we get anything decent from reading a bunch of random books in

Requires an additional folder “decades” in the root directory

```
[11]: # Get all filenames
files = [f for f in listdir(file_path) if isfile(join(file_path, f))]

# Do only subset
## Is done for 5000 files already, so set down to 20 to increase performance.
→ 5000 books are currently stored in the file
files = files[0:20]

counter = 0
for file in files:
    counter = counter + 1
    # Read in basic information from file
    title, author, date, year, language, encoding, content_lines =
→ read_file(file_path + "/" + file)
    #line_count = len(content_lines)
```



```

decade = math.floor(year / 10) * 10
decade_file = "decades/" + str(decade) + ".txt"
content_all = " ".join(content_lines)

if os.path.exists(decade_file):
    append_write = 'a' # append if already exists
else:
    append_write = 'w' # make a new file if not

fileWriter = open(decade_file,append_write)
fileWriter.write(content_all + '\n')
fileWriter.close()

```

#### 2.4.1 Read in from the decades files, and see the distributions

```

[12]: # Get all filenames
files = [f for f in listdir("decades") if isfile(join("decades", f))]
print(files)
files.sort(reverse=True)

col_names = []
col_names.append("Word")

tables = []

for file_name in files:
    print(file_name)

    file = open("decades/" + file_name, encoding="ISO-8859-1")
    file_content = file.read()

    # Split into words (and do various cleaning)
    all_text_lower = file_content.lower()
    words = re.findall(r'(\b[A-Za-z][a-z]{2,9}\b)', all_text_lower)

    # First analysis, but should do something proper
    word_frequencies_table, unique_word_count = get_word_frequencies(words)
    tables.append(word_frequencies_table)
    col_names.append(file_name)

```

```

['00.txt', '0.txt', '2010.txt', '2000.txt', '2020.txt', '1990.txt']
2020.txt
2010.txt
2000.txt
1990.txt
00.txt

```

0.txt

### 2.4.2 Preliminary Conclusion

We see that even though the books are quite old, no decade prior to 1990s is found.

This is when we found out that the “year” that’s registered in the dataset is the upload-date.

Haven gotten this far, we however decided to see if we could find a pattern in this

### 2.4.3 Compare ranking between upload-decades

```
[13]: list_count= []
list_freq = []
list_rank = []

for df in tables:
    #list_count.append(df[['Word', 'count']])
    #list_freq.append(df[['Word', 'freq']])
    list_rank.append(df[['Word', 'rank']])

#df_count = reduce(lambda left, right: pd.merge(left, right, on="Word",
    ↳how='outer'), list_count)
#df_count.columns = col_names

#df_freq = reduce(lambda left, right: pd.merge(left, right, on="Word",
    ↳how='outer'), list_freq)
#df_freq.columns = col_names

df_rank = reduce(lambda left, right: pd.merge(left, right, on="Word",
    ↳how='outer'), list_rank)
df_rank.columns = col_names
```

```
[14]: df_rank.head(100)
```

```
[14]:
```

	Word	2020.txt	2010.txt	2000.txt	1990.txt	00.txt	0.txt
0	the	1.0	1.0	1.0	1.0	1.0	1.0
1	and	2.0	2.0	2.0	2.0	2.0	2.0
2	that	3.0	3.0	3.0	3.0	4.0	3.0
3	was	4.0	4.0	4.0	4.0	23.0	5.0
4	you	5.0	10.0	8.0	5.0	163.0	19.0
5	with	6.0	6.0	6.0	7.0	3.0	6.0
6	for	7.0	7.0	7.0	11.0	13.0	8.0
7	his	8.0	5.0	5.0	6.0	5.0	4.0
8	not	9.0	8.0	11.0	12.0	11.0	9.0
9	had	10.0	9.0	9.0	10.0	40.0	14.0
10	but	11.0	11.0	10.0	13.0	72.0	10.0
11	which	12.0	12.0	15.0	22.0	82.0	7.0

12	they	13.0	17.0	14.0	16.0	58.0	21.0
13	from	14.0	14.0	18.0	21.0	25.0	15.0
14	were	15.0	20.0	21.0	20.0	81.0	22.0
15	have	16.0	16.0	16.0	18.0	22.0	11.0
16	this	17.0	13.0	17.0	15.0	8.0	13.0
17	are	18.0	18.0	24.0	27.0	37.0	16.0
18	she	19.0	21.0	13.0	9.0	65.0	47.0
19	all	20.0	19.0	20.0	17.0	15.0	12.0
20	their	21.0	24.0	25.0	26.0	21.0	30.0
21	him	22.0	22.0	19.0	14.0	28.0	20.0
22	her	23.0	15.0	12.0	8.0	20.0	33.0
23	its	24.0	40.0	55.0	98.0	68.0	67.0
24	one	25.0	23.0	22.0	25.0	17.0	28.0
25	there	26.0	25.0	23.0	23.0	37.0	32.0
26	them	27.0	28.0	28.0	32.0	55.0	38.0
27	what	28.0	33.0	32.0	24.0	56.0	24.0
28	has	29.0	36.0	46.0	46.0	176.0	40.0
29	been	30.0	29.0	31.0	33.0	30.0	23.0
30	will	31.0	31.0	33.0	37.0	85.0	27.0
31	would	32.0	30.0	29.0	31.0	194.0	25.0
32	said	33.0	32.0	26.0	19.0	61.0	141.0
33	when	34.0	27.0	27.0	28.0	105.0	34.0
34	more	35.0	34.0	36.0	42.0	190.0	26.0
35	who	36.0	26.0	30.0	30.0	27.0	18.0
36	into	37.0	37.0	37.0	39.0	97.0	63.0
37	out	38.0	35.0	34.0	29.0	62.0	77.0
38	then	39.0	38.0	35.0	35.0	47.0	50.0
39	other	40.0	44.0	56.0	58.0	88.0	66.0
40	men	41.0	72.0	78.0	81.0	108.0	53.0
41	only	42.0	48.0	53.0	60.0	192.0	71.0
42	can	43.0	59.0	48.0	45.0	98.0	54.0
43	upon	44.0	51.0	57.0	80.0	121.0	89.0
44	our	45.0	55.0	49.0	91.0	149.0	43.0
45	than	46.0	47.0	50.0	64.0	201.0	31.0
46	now	47.0	43.0	38.0	38.0	134.0	42.0
47	time	48.0	42.0	43.0	49.0	96.0	57.0
48	power	49.0	241.0	257.0	327.0	204.0	111.0
49	great	50.0	62.0	63.0	79.0	64.0	37.0
50	these	51.0	49.0	67.0	87.0	48.0	60.0
51	government	52.0	379.0	440.0	534.0	NaN	146.0
52	man	53.0	41.0	41.0	34.0	222.0	29.0
53	over	54.0	67.0	64.0	62.0	92.0	139.0
54	could	55.0	45.0	40.0	41.0	199.0	64.0
55	very	56.0	46.0	42.0	50.0	159.0	65.0
56	your	57.0	57.0	45.0	36.0	201.0	55.0
57	first	58.0	64.0	76.0	77.0	219.0	81.0
58	society	59.0	579.0	603.0	535.0	NaN	231.0

59	two	60.0	52.0	59.0	70.0	124.0	100.0
60	made	61.0	61.0	62.0	83.0	174.0	82.0
61	such	62.0	63.0	75.0	78.0	146.0	45.0
62	about	63.0	53.0	44.0	40.0	99.0	113.0
63	some	64.0	39.0	39.0	44.0	70.0	49.0
64	any	65.0	54.0	54.0	55.0	130.0	49.0
65	did	66.0	60.0	52.0	54.0	154.0	99.0
66	know	67.0	83.0	74.0	48.0	159.0	98.0
67	pendleton	68.0	3542.0	3792.0	773.0	NaN	NaN
68	same	69.0	97.0	125.0	118.0	203.0	112.0
69	well	70.0	68.0	58.0	52.0	54.0	58.0
70	under	71.0	103.0	121.0	153.0	129.0	103.0
71	may	72.0	50.0	70.0	86.0	162.0	36.0
72	general	73.0	196.0	225.0	374.0	254.0	195.0
73	before	74.0	66.0	65.0	59.0	101.0	84.0
74	most	75.0	80.0	88.0	120.0	243.0	46.0
75	even	76.0	88.0	101.0	125.0	77.0	78.0
76	much	77.0	77.0	77.0	76.0	228.0	73.0
77	like	78.0	65.0	51.0	47.0	10.0	74.0
78	stephanie	79.0	NaN	3771.0	NaN	NaN	NaN
79	lorraine	80.0	3413.0	3652.0	766.0	NaN	383.0
80	those	81.0	79.0	91.0	138.0	26.0	44.0
81	down	82.0	76.0	68.0	63.0	63.0	126.0
82	back	83.0	98.0	86.0	85.0	165.0	217.0
83	came	84.0	93.0	85.0	74.0	172.0	236.0
84	see	85.0	69.0	60.0	51.0	150.0	88.0
85	how	86.0	84.0	73.0	57.0	152.0	56.0
86	way	87.0	87.0	82.0	69.0	118.0	145.0
87	think	88.0	125.0	103.0	82.0	227.0	114.0
88	little	89.0	56.0	47.0	43.0	229.0	80.0
89	without	90.0	107.0	107.0	107.0	105.0	79.0
90	here	91.0	82.0	81.0	68.0	89.0	93.0
91	against	92.0	134.0	144.0	171.0	220.0	92.0
92	people	93.0	111.0	116.0	121.0	203.0	124.0
93	after	94.0	58.0	61.0	53.0	103.0	95.0
94	must	95.0	71.0	79.0	88.0	213.0	69.0
95	don	95.0	148.0	102.0	56.0	NaN	274.0
96	where	96.0	75.0	71.0	84.0	151.0	114.0
97	never	97.0	90.0	83.0	75.0	172.0	83.0
98	own	98.0	92.0	92.0	96.0	111.0	68.0
99	right	99.0	133.0	135.0	117.0	236.0	161.0

## 2.5 Trying to fit models to predict

### 2.5.1 Read in files

```
[15]: file_contents = []
      targets = []

      files = [f for f in listdir(file_path) if isfile(join(file_path, f))]
      files = list(filter(lambda file: file[0].isdigit(), files))
      random.shuffle(files)

      targets_=['70','80','90','00','10']
      iter_ = 0

      for f in files[:120]:
          file = open("processedData/" + f, encoding="ISO-8859-1")
          file_contents.append(file.read())
          iter_ = iter_+1
          targets.append(targets_[iter_%5])
```

### 2.5.2 Train models

```
[16]: pipeline = Pipeline([
      ('vect', CountVectorizer()),
      ('tfidf', TfidfTransformer()),
      ('kbest', SelectKBest(chi2, k=100)),
      ('nb', MultinomialNB()),
  ])

  parameters = {
      #'vect__max_df': [1.0],
      # 'vect__max_features': (None, 5000, 10000, 50000),
      #'vect__ngram_range': ((1, 1), (1, 2)), # unigrams or bigrams
      # 'tfidf__use_idf': (True, False),
      # 'tfidf__norm': ('l1', 'l2'),
      #'clf__max_iter': (20),
      #'clf__alpha': (0.00001),
      #'clf__penalty': ('l2'),
      # 'clf__max_iter': (10, 50, 80),
  }

  grid_search = GridSearchCV(pipeline, parameters, verbose=1)

  grid_search.fit(file_contents, targets)
  best_parameters = grid_search.best_estimator_.get_params()

  for param_name in sorted(parameters.keys()):
      print("\t%s: %r" % (param_name, best_parameters[param_name]))
```

Fitting 5 folds for each of 1 candidates, totalling 5 fits

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.  
[Parallel(n_jobs=1)]: Done 5 out of 5 | elapsed: 36.4s finished
```

## 2.6 Realisation and conclusion

At this point, we came to the conclusion that “year” in the Gutenberg dataset shows when the data **was published** to the project, and not the release date of the book.

We searched for possible solutions to get the years for book publications, but were unable to find any free API that we could link to our current dataset.

We therefore went on a search for other datasets, and to remake our hypothesis entirely. Thus, this part ended in a blind spot. However science is not only about the results, but also about the discoveries along the way, and therefore it is added into this file.

## 3 Studying language change in Icelandic parliamentary speeches

Our task involves research into **language change over the past 100 years**. Additionally we have been tasked with working out factors that influence language change.

Another proposed research question could have been focused on figuring out which languages are going extinct. This particular task has been found out to be near impossible to answer given the available data. It is estimated to be very hard to come up with data that capture the amount of speakers for a large enough ranges of combinations of language and year. Furthermore, any data that are available are likely to apply a different definition of “speaker” (sometimes including second/third... language speakers, sometimes not) and is also likely to contain politically motivated noise.

### 3.1 Introduction

Therefore, we decided to search for English language corpora containing a wide array of text documents collected over the past century for predefined dialects of English and genre of text (movie, articles, books, ...). This surprisingly turned out to be a complex endeavour as all high quality corpora were available only for a big price tag.

We also looked into the material provided by the Guttenberg Project [Link](#). This turned out to be promising at first sight as it appears that there is a lot of recently published material. However release date of these documents does not match the year when the documents were actually written and soon enough we figured out that all material is from before 1923. This obviously did not allow us to look much into language change of the 20th and 21st century.

*Gerlach, M., & Font-Clos, F. (2020). A standardized Project Gutenberg corpus for statistical analysis of natural language and quantitative linguistics. Entropy, 22(1), 126.*

Theoretically one could obtain books from after 1923 and include them into the analysis. But one would quickly run into copyright/licensing issues here.

Obtaining the content of these books and preprocessing them for the purposes of data analysis turned out to be quite cumbersome as well. Look at Gunnar’s notebooks (first draft [here](#), second draft [here](#)) for the details.

Finally we turned to looking for non-English corpora and **found an annotated corpus including pre-factured lemmatization of Icelandic parliamentary speeches from 1911 until 2018:**

*Steingrímsson, Steinþór, Sigrún Helgadóttir, Eiríkur Rögnvaldsson, Starkaður Barkarson and Jón Guðnason. 2018. Risamálheild: A Very Large Icelandic Text Corpus. Proceedings of LREC 2018, pp. 4361-4366. Myazaki, Japan.*

## 3.2 The task

In the line with our goal of analyzing the change in language over the past 100 years, we decided to train different models and assess their ability to predict whether an speech held in the Icelandic Parliament belongs to a particular decade. In the end, this is a **document classification task** in which the input is a large set of parliament speeches and the target/class is the decade in which the speeches were held.

A good performance of our proposed classifiers may support the idea that Icelandic has envolved in the years. However, the fact that the models would perform well is not enough to assert that the language has changed. It could be that what has actually changed are the topics or even the way of documenting the speeches. Anyway, for us it was really exciting to check whether we are able to **fit a model that predicts reasonably well the decade of an speech by only using the speech itself.**

## 3.3 Setup

In this section, we provide the **setup for a successful** implementation (or replication) of our experiment within this Jupyter Notebook.

### 3.3.1 Load required libraries

The following libraries are used during the next sections and therefore need to be imported.

```
[17]: import pandas as pd
import numpy as np
import xml.etree.ElementTree as ET
import glob
from nltk.probability import FreqDist
import random
from functools import reduce
from nltk import ngrams
# Used for building models for classifying:
from pprint import pprint
from time import time
import logging
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.linear_model import SGDClassifier, LogisticRegression
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import make_scorer, accuracy_score
from sklearn.model_selection import GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.svm import SVC
```

```

from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.base import BaseEstimator
from sklearn.metrics import accuracy_score
from sklearn.metrics import f1_score
from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import confusion_matrix
from sklearn.feature_selection import SelectKBest, chi2
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
import matplotlib.pyplot as plt
from gensim.models import Word2Vec
from gensim.models import Doc2Vec
from gensim.models.doc2vec import TaggedDocument
from sklearn.metrics import plot_confusion_matrix

```

```

[18]: #needed afterwards too
namespace = "{http://www.tei-c.org/ns/1.0}"

```

### 3.3.2 Get the data

Data can be downloaded from here: <http://www.malfong.is/index.php?dlid=81&lang=en>. However, we provided already in our submission file the specifications on how to get the data of our assignment.

Then extract zip folder such that a folder labelled CC\_BY shows up in the parent folder of this notebook. *Test:* `ls ../CC_BY/althingi` should work when run from `.../IcelandicParliamentSpeeches.ipynb`.

### 3.3.3 Preprocessing helpers

The data are available as XML. The text has already been preprocessed to be separated into paragraphs, sentences and words. Furthermore each word tag also includes a `lemma` attribute relating inflected/declensed forms of words to its lemma. This has been done by the authors of the original paper using Machine Learning approaches.

Given a relative path to a file, pull out a list with all the words. This can be achieved by looking for all tags of type `w`, additionally also retrieve the lemma for each word.

We will discard all sentences of length 3 or smaller to remove noise and to avoid that our models are able to detect year of speech just based on some short introductory/outro phrases. Furthermore the raw data appear to contain plenty of elements tagged as words that comprise of just a single letter followed by a dot. These will be removed here as well.

*Pitfall:* The namespace from above must be included when parsing out content from these XML files based on tag names.



In this kind of preprocessing we lose information about sentence boundaries as all punctuation items from the raw data are dropped.

```
[19]: def extract_words(path):
    xml_tree = ET.parse(open(path, 'r', encoding="utf8"))
    words = []
    lemmata = []

    for sentence in xml_tree.getroot().iter('{s'.format(namespace)):
        words_in_sent = sentence.findall('{w'.format(namespace))
        if len(words_in_sent) > 2:
            for word in sentence.findall('{w'.format(namespace)):
                if not word.text.endswith('.'):
                    words.append(word.text)
                    lemmata.append(word.attrib['lemma'])

    return words, lemmata
```

Extract content of files separated into sentences, note that all stop items are wrapped in a p tag in the original documents and are not included here.

Also note that some further pre-processing could be done here to exclude items such as numbers, percentages, names, abbreviations, etc. In the original documents these are also assigned to be words:

```
[20]: def extract_sentences(path, lemma=False):
    xml_tree = ET.parse(open(path, 'r', encoding="utf8"))
    sentences = []

    for sentence in xml_tree.getroot().iter('{s'.format(namespace)):
        sentence_cur = []
        words_in_sent = sentence.findall('{w'.format(namespace))

        if len(words_in_sent) > 2:
            for word in sentence.findall('{w'.format(namespace)):
                if not word.text.endswith('.'):
                    if lemma:
                        sentence_cur.append(word.attrib['lemma'])
                    else:
                        sentence_cur.append(word.text)

            sentences.append(sentence_cur)

    return sentences
```

Retrieve a random selection of k file names from the entire corpus. The files must be of type xml. This method does not load the entire corpus into memory and allows you to work with smaller selections for test purposes. This method samples only from the althingi folder so far:

```
[21]: def get_random_sample(k):
        files = [filename for filename in glob.iglob('../CC_BY/althingi/**/*.xml',
                                                    recursive=True)]

        return random.sample(files, k)
```

```
[22]: files = [filename for filename in glob.iglob('../CC_BY/althingi/**/*.xml',
↪recursive=True)]
        #print(files)
```

Do the same as above but choose k files only from a given year (range: 1911-2017)

```
[23]: def get_files_for_year(year, k = None):
        files = [filename for filename in glob.iglob('../CC_BY/althingi/{}/'.
↪format(year) + '**/*.xml',
                                                    recursive=True)]

        if k == None:
            newK = len(files)
            res = files
        else:
            newK = k
            res = random.sample(files, min(len(files), k))
        if len(files) != 0:
            percentage = 100*newK/len(files)
        else:
            percentage = 0
        print("For year " + str(year) + ": Fetching " + str(newK) + " samples out of "
↪" + str(len(files)) + " (~" + str(percentage) + "%)")
        return res
```

### 3.4 Preliminary Data Analysis

In this section, we perform a preliminary data analysis to get a better insight of our data.

#### 3.4.1 Zipf's Law


First using frequency distributions of the Natural Language ToolKit (NLTK) to look into whether or not we can confirm [Zipf's Law](#) based on the data we have.

Note that the analysis is done based on 15 randomly selected files from the entire corpus at this point:

```
[24]: words = []

        for file in get_random_sample(15):
            words.extend(extract_words(file)[1])

        fq = FreqDist(word.lower() for word in words)
        fq.plot(25, cumulative=False)
```




[24]: <AxesSubplot:xlabel='Samples', ylabel='Counts'>

Visualizing the same data but with using the logarithm of the occurrences, this should ideally obtain a straight line:

```
[25]: freq_df = pd.DataFrame.from_dict(fq, orient='index', columns=['word_occur'])
freq_df.sort_values(by='word_occur', inplace=True, ascending=False)
freq_df.word_occur = np.log2(freq_df['word_occur'])
freq_df.head(25).plot(kind='line')
```

[25]: <AxesSubplot:>



### 3.4.2 Disappearing words / new words

Here is a description

```
[26]: words_1914 = []
      words_2014 = []

      for file in get_files_for_year(1914, 25):
          words_1914.extend(extract_words(file)[1])

      for file in get_files_for_year(2014, 25):
          words_2014.extend(extract_words(file)[1])
```

For year 1914: Fetching 25 samples out of 1306 (~1.9142419601837672%)

For year 2014: Fetching 25 samples out of 12404 (~0.2015478877781361%)

### 3.4.3 Development of average sentence length

This is just one possible metric for the development/analysis of language complexity. There is so much more you could come up with here.

Obviously our choice to discard very short sentences in the preprocessing step has an impact on the values here:

```
[27]: def avg_sentence_length_year(year, k):
      sentence_len = []
      for file in get_files_for_year(year, k):
```

```

    sentences = extract_sentences(file)
    sentence_len.extend([len(s) for s in sentences])

    return reduce(lambda a, b: a + b, sentence_len) / len(sentence_len)

sentence_len_years = []

for year in range(1924, 2018):
    sentence_len_years.append(avg_sentence_length_year(year, 20))

avg_df = pd.DataFrame(sentence_len_years, index=range(1924, 2018),
    ↳columns=['avg_sent_len'])
avg_df.plot(kind='line')

```

```

For year 1924: Fetching 20 samples out of 41 (~48.78048780487805%)
For year 1925: Fetching 20 samples out of 23 (~86.95652173913044%)
For year 1926: Fetching 20 samples out of 37 (~54.054054054054056%)
For year 1927: Fetching 20 samples out of 59 (~33.898305084745765%)
For year 1928: Fetching 20 samples out of 21 (~95.23809523809524%)
For year 1929: Fetching 20 samples out of 96 (~20.833333333333332%)
For year 1930: Fetching 20 samples out of 30 (~66.66666666666667%)
For year 1931: Fetching 20 samples out of 25 (~80.0%)
For year 1932: Fetching 20 samples out of 28 (~71.42857142857143%)
For year 1933: Fetching 20 samples out of 52 (~38.46153846153846%)
For year 1934: Fetching 20 samples out of 20 (~100.0%)
For year 1935: Fetching 20 samples out of 24 (~83.33333333333333%)
For year 1936: Fetching 20 samples out of 25 (~80.0%)
For year 1937: Fetching 20 samples out of 1381 (~1.448225923244026%)
For year 1938: Fetching 20 samples out of 1676 (~1.1933174224343674%)
For year 1939: Fetching 20 samples out of 1632 (~1.2254901960784315%)
For year 1940: Fetching 20 samples out of 1458 (~1.3717421124828533%)
For year 1941: Fetching 20 samples out of 2066 (~0.968054211035818%)
For year 1942: Fetching 20 samples out of 2357 (~0.8485362749257531%)
For year 1943: Fetching 20 samples out of 3960 (~0.5050505050505051%)
For year 1944: Fetching 20 samples out of 1072 (~1.8656716417910448%)
For year 1945: Fetching 20 samples out of 1859 (~1.0758472296933836%)
For year 1946: Fetching 20 samples out of 2789 (~0.7171029042667623%)
For year 1947: Fetching 20 samples out of 2838 (~0.704721634954193%)
For year 1948: Fetching 20 samples out of 2262 (~0.8841732979664014%)
For year 1949: Fetching 20 samples out of 2661 (~0.7515971439308531%)
For year 1950: Fetching 20 samples out of 3017 (~0.6629101756711966%)
For year 1951: Fetching 20 samples out of 2453 (~0.8153281695882593%)
For year 1952: Fetching 20 samples out of 1689 (~1.1841326228537596%)
For year 1953: Fetching 20 samples out of 1415 (~1.4134275618374559%)
For year 1954: Fetching 20 samples out of 1546 (~1.2936610608020698%)
For year 1955: Fetching 20 samples out of 1433 (~1.3956734124214933%)
For year 1956: Fetching 20 samples out of 1213 (~1.6488046166529267%)

```


For year 1957: Fetching 20 samples out of 1763 (~1.1344299489506522%)  
 For year 1958: Fetching 20 samples out of 1169 (~1.7108639863130881%)  
 For year 1959: Fetching 20 samples out of 1326 (~1.5082956259426847%)  
 For year 1960: Fetching 20 samples out of 1862 (~1.0741138560687433%)  
 For year 1961: Fetching 20 samples out of 1409 (~1.4194464158978%)  
 For year 1962: Fetching 20 samples out of 1475 (~1.3559322033898304%)  
 For year 1963: Fetching 20 samples out of 1376 (~1.4534883720930232%)  
 For year 1964: Fetching 20 samples out of 1677 (~1.1926058437686344%)  
 For year 1965: Fetching 20 samples out of 1520 (~1.3157894736842106%)  
 For year 1966: Fetching 20 samples out of 1404 (~1.4245014245014245%)  
 For year 1967: Fetching 20 samples out of 1274 (~1.5698587127158556%)  
 For year 1968: Fetching 20 samples out of 1726 (~1.1587485515643106%)  
 For year 1969: Fetching 20 samples out of 1740 (~1.1494252873563218%)  
 For year 1970: Fetching 20 samples out of 1921 (~1.041124414367517%)  
 For year 1971: Fetching 20 samples out of 2071 (~0.9657170449058425%)  
 For year 1972: Fetching 20 samples out of 2368 (~0.8445945945945946%)  
 For year 1973: Fetching 20 samples out of 2329 (~0.8587376556462001%)  
 For year 1974: Fetching 20 samples out of 2350 (~0.851063829787234%)  
 For year 1975: Fetching 20 samples out of 2430 (~0.823045267489712%)  
 For year 1976: Fetching 20 samples out of 2555 (~0.7827788649706457%)  
 For year 1977: Fetching 20 samples out of 2105 (~0.9501187648456056%)  
 For year 1978: Fetching 20 samples out of 2641 (~0.7572889057175313%)  
 For year 1979: Fetching 20 samples out of 2165 (~0.9237875288683602%)  
 For year 1980: Fetching 20 samples out of 3406 (~0.5871990604815033%)  
 For year 1981: Fetching 20 samples out of 3491 (~0.5729017473503294%)  
 For year 1982: Fetching 20 samples out of 2894 (~0.691085003455425%)  
 For year 1983: Fetching 20 samples out of 2651 (~0.7544322897019993%)  
 For year 1984: Fetching 20 samples out of 4171 (~0.4795013186286262%)  
 For year 1985: Fetching 20 samples out of 4709 (~0.42471862391165854%)  
 For year 1986: Fetching 20 samples out of 3432 (~0.5827505827505828%)  
 For year 1987: Fetching 20 samples out of 3212 (~0.6226650062266501%)  
 For year 1988: Fetching 20 samples out of 4198 (~0.47641734159123394%)  
 For year 1989: Fetching 20 samples out of 5085 (~0.39331366764995085%)  
 For year 1990: Fetching 20 samples out of 4662 (~0.429000429000429%)  
 For year 1991: Fetching 20 samples out of 4747 (~0.4213187276174426%)  
 For year 1992: Fetching 20 samples out of 8925 (~0.22408963585434175%)  
 For year 1993: Fetching 20 samples out of 7412 (~0.26983270372369134%)  
 For year 1994: Fetching 20 samples out of 8187 (~0.2442897276169537%)  
 For year 1995: Fetching 20 samples out of 5129 (~0.38993955936829794%)  
 For year 1996: Fetching 20 samples out of 7184 (~0.27839643652561247%)  
 For year 1997: Fetching 20 samples out of 6960 (~0.28735632183908044%)  
 For year 1998: Fetching 20 samples out of 7393 (~0.27052617340727714%)  
 For year 1999: Fetching 20 samples out of 6056 (~0.33025099075297226%)  
 For year 2000: Fetching 20 samples out of 7466 (~0.2678810608090008%)  
 For year 2001: Fetching 20 samples out of 8210 (~0.243605359317905%)  
 For year 2002: Fetching 20 samples out of 8061 (~0.24810817516437167%)  
 For year 2003: Fetching 20 samples out of 5872 (~0.3405994550408719%)  
 For year 2004: Fetching 20 samples out of 9466 (~0.21128248468201985%)

```

For year 2005: Fetching 20 samples out of 8269 (~0.24186721489902044%)
For year 2006: Fetching 20 samples out of 8810 (~0.22701475595913734%)
For year 2007: Fetching 20 samples out of 7863 (~0.2543558438255119%)
For year 2008: Fetching 20 samples out of 8764 (~0.22820629849383842%)
For year 2009: Fetching 20 samples out of 17262 (~0.11586142973004288%)
For year 2010: Fetching 20 samples out of 11089 (~0.18035891423933628%)
For year 2011: Fetching 20 samples out of 13957 (~0.14329727018700295%)
For year 2012: Fetching 20 samples out of 16356 (~0.12227928588897041%)
For year 2013: Fetching 20 samples out of 10240 (~0.1953125%)
For year 2014: Fetching 20 samples out of 12404 (~0.16123831022250887%)
For year 2015: Fetching 20 samples out of 18052 (~0.11079104808331487%)
For year 2016: Fetching 20 samples out of 8165 (~0.2449479485609308%)
For year 2017: Fetching 20 samples out of 7270 (~0.2751031636863824%)

```

[27]: <AxesSubplot:>



### 3.4.4 n-grams

Here is a description

```

[28]: def most_common_ngrams(n, top_k, sample):
        file_contents = []

        for file in get_random_sample(sample):
            file_contents.extend(extract_words(file)[1])


```

```

fq_ngr = FreqDist(ngrams(file_contents, n))
fq_ngr.plot(top_k, cumulative=False)

most_common_ngrams(n=3, top_k=25, sample=100)

```



### 3.5 Building model for classifying speeches

This section is the main core of the modelling task of this assignment. It is organized as follows:

- Firstly, the data is splitted into **training and test sets** according to the criteria that in our opinion fits better to the nature of data and task.
- Secondly, we describe the **methods** that are going to be used for **feature extraction** from our documents.
- After that, we describe the **classifiers chosen** to be trained and why they were selected.
- Then, we **train 7 models** combining the feature extraction techniques described and the classifiers selected. This is done through a cross-validated grid search in which many hyperparameters are combined. The goal of this search is to find the **best hyperparameter**



combination of each of the 7 models.

- Finally, we **compare the results from the training within and between the models**.
- Evaluation on test data will be performed in the next section.

### 3.5.1 Constructing training and test data

Our whole dataset contains **380285 speeches** hold in the Icelandic parliament from 1911 to 2017. In order to perform our train-test split, we took into account the following considerations:

- Documents are **classified in directories by year and month instead of decade**.
- Decades (classes) are highly **unbalanced**. There are much more documents from later decades as from the earlier ones. As an example, 1912 has only 14 documents while 2011 has 13957. This may introduce bias in the training of the models if not dealt.

To solve the first problem, we use the help function `get_files_for_year()` created above, which takes `n` documents from an specified year. After that, for each of the documents, the year is substituted by the decade as shown in the next two sections. This can be done iteratively through a list of years. In this way, **we obtain a dataset with a bunch of corpora labelled by decade**.

To solve the problem of unbalance within classes, we **limit the number of documents to be extracted from each year to 200 for the training set**. This way, we ensure that there will not be too big differences within the number of documents sampled within the years (maximum of 200 vs minimum of 14) and neither within the decades. We choose 200 since we consider it to be a good balance for **undersampling the majority classes but not losing as much information as we would keep it to minimum of 14**.

Note that, in order to make the runtimes of our notebook shorter (feasible) we **skip intermediate decades from our classification task**. This would simulate that there were not speeches held in some decades. We like to imagine it as weird regime which combines a decade of democracy followed by a decade of dictatorship. In summary:

- 1910s, 1930s, 1950s, 1970s, 1990s, and 2010s are considered.
- Whereas 1920s, 1940s, 1960s, 1980s, and 2000s are discarded.

We will perform a train/test split of the approximate proportion 80/20. We will see why it will be approximate in the next two sections.

**Train data** 8 years out of the 10 years that form a decade are chosen for each of the 6 decades considered for the train set. **The other 2 are left for the test set**. The selection of the years was completely random. For each of the decades a maximum of 1600 documents are chosen. However, this will not be equal for all the decades, since, as explained above, not all the years have at least 200 documents.

Note that for the last decade, we just have documents until 2017. The split will be 6 years (train) vs 1 (test) in this case. Same applies for first decade (in this case, 7 vs. 2).

```
[29]: #set seed for reproducibility
      random.seed(123)

      file_contents = []
      targets = []
```

```

for year in [1911, 1912, 1914, 1915, 1916, 1918, 1919,
            1920, 1921, 1922, 1924, 1925, 1926, 1928, 1929,
            1931, 1932, 1933, 1934, 1935, 1936, 1937, 1938,
            1941, 1942, 1943, 1944, 1945, 1946, 1947, 1948,
            1951, 1952, 1953, 1955, 1956, 1957, 1958, 1959,
            1961, 1962, 1963, 1965, 1966, 1967, 1968, 1969,
            1970, 1971, 1972, 1973, 1974, 1975, 1978, 1979,
            1980, 1981, 1982, 1983, 1984, 1985, 1988, 1989,
            1990, 1991, 1992, 1993, 1995, 1996, 1997, 1999,
            2000, 2001, 2002, 2003, 2005, 2006, 2007, 2009,
            2010, 2011, 2012, 2013, 2014, 2016, 2017]:
    for file in get_files_for_year(year, 200):
        file_contents.append(extract_words(file)[1])
        targets.append(year - year%10)

```

```

For year 1911: Fetching 200 samples out of 125 (~160.0%)
For year 1912: Fetching 200 samples out of 14 (~1428.5714285714287%)
For year 1914: Fetching 200 samples out of 1306 (~15.313935681470138%)
For year 1915: Fetching 200 samples out of 1383 (~14.461315979754158%)
For year 1916: Fetching 200 samples out of 0 (~0%)
For year 1918: Fetching 200 samples out of 0 (~0%)
For year 1919: Fetching 200 samples out of 0 (~0%)
For year 1920: Fetching 200 samples out of 0 (~0%)
For year 1921: Fetching 200 samples out of 0 (~0%)
For year 1922: Fetching 200 samples out of 0 (~0%)
For year 1924: Fetching 200 samples out of 41 (~487.8048780487805%)
For year 1925: Fetching 200 samples out of 23 (~869.5652173913044%)
For year 1926: Fetching 200 samples out of 37 (~540.5405405405405%)
For year 1928: Fetching 200 samples out of 21 (~952.3809523809524%)
For year 1929: Fetching 200 samples out of 96 (~208.33333333333334%)
For year 1931: Fetching 200 samples out of 25 (~800.0%)
For year 1932: Fetching 200 samples out of 28 (~714.2857142857143%)
For year 1933: Fetching 200 samples out of 52 (~384.61538461538464%)
For year 1934: Fetching 200 samples out of 20 (~1000.0%)
For year 1935: Fetching 200 samples out of 24 (~833.3333333333334%)
For year 1936: Fetching 200 samples out of 25 (~800.0%)
For year 1937: Fetching 200 samples out of 1381 (~14.48225923244026%)
For year 1938: Fetching 200 samples out of 1676 (~11.933174224343675%)
For year 1941: Fetching 200 samples out of 2066 (~9.68054211035818%)
For year 1942: Fetching 200 samples out of 2357 (~8.485362749257531%)
For year 1943: Fetching 200 samples out of 3960 (~5.05050505050505%)
For year 1944: Fetching 200 samples out of 1072 (~18.65671641791045%)
For year 1945: Fetching 200 samples out of 1859 (~10.758472296933835%)
For year 1946: Fetching 200 samples out of 2789 (~7.171029042667623%)
For year 1947: Fetching 200 samples out of 2838 (~7.047216349541931%)

```

For year 1948: Fetching 200 samples out of 2262 (~8.841732979664014%)  
 For year 1951: Fetching 200 samples out of 2453 (~8.153281695882592%)  
 For year 1952: Fetching 200 samples out of 1689 (~11.841326228537596%)  
 For year 1953: Fetching 200 samples out of 1415 (~14.134275618374557%)  
 For year 1955: Fetching 200 samples out of 1433 (~13.956734124214934%)  
 For year 1956: Fetching 200 samples out of 1213 (~16.488046166529266%)  
 For year 1957: Fetching 200 samples out of 1763 (~11.344299489506524%)  
 For year 1958: Fetching 200 samples out of 1169 (~17.108639863130882%)  
 For year 1959: Fetching 200 samples out of 1326 (~15.082956259426847%)  
 For year 1961: Fetching 200 samples out of 1409 (~14.194464158977999%)  
 For year 1962: Fetching 200 samples out of 1475 (~13.559322033898304%)  
 For year 1963: Fetching 200 samples out of 1376 (~14.534883720930232%)  
 For year 1965: Fetching 200 samples out of 1520 (~13.157894736842104%)  
 For year 1966: Fetching 200 samples out of 1404 (~14.245014245014245%)  
 For year 1967: Fetching 200 samples out of 1274 (~15.698587127158556%)  
 For year 1968: Fetching 200 samples out of 1726 (~11.587485515643106%)  
 For year 1969: Fetching 200 samples out of 1740 (~11.494252873563218%)  
 For year 1970: Fetching 200 samples out of 1921 (~10.41124414367517%)  
 For year 1971: Fetching 200 samples out of 2071 (~9.657170449058427%)  
 For year 1972: Fetching 200 samples out of 2368 (~8.445945945945946%)  
 For year 1973: Fetching 200 samples out of 2329 (~8.587376556462%)  
 For year 1974: Fetching 200 samples out of 2350 (~8.51063829787234%)  
 For year 1975: Fetching 200 samples out of 2430 (~8.23045267489712%)  
 For year 1978: Fetching 200 samples out of 2641 (~7.5728890571753125%)  
 For year 1979: Fetching 200 samples out of 2165 (~9.237875288683602%)  
 For year 1980: Fetching 200 samples out of 3406 (~5.871990604815032%)  
 For year 1981: Fetching 200 samples out of 3491 (~5.729017473503294%)  
 For year 1982: Fetching 200 samples out of 2894 (~6.91085003455425%)  
 For year 1983: Fetching 200 samples out of 2651 (~7.544322897019993%)  
 For year 1984: Fetching 200 samples out of 4171 (~4.795013186286262%)  
 For year 1985: Fetching 200 samples out of 4709 (~4.247186239116585%)  
 For year 1988: Fetching 200 samples out of 4198 (~4.764173415912339%)  
 For year 1989: Fetching 200 samples out of 5085 (~3.933136676499508%)  
 For year 1990: Fetching 200 samples out of 4662 (~4.29000429000429%)  
 For year 1991: Fetching 200 samples out of 4747 (~4.213187276174426%)  
 For year 1992: Fetching 200 samples out of 8925 (~2.2408963585434174%)  
 For year 1993: Fetching 200 samples out of 7412 (~2.698327037236913%)  
 For year 1995: Fetching 200 samples out of 5129 (~3.899395593682979%)  
 For year 1996: Fetching 200 samples out of 7184 (~2.7839643652561246%)  
 For year 1997: Fetching 200 samples out of 6960 (~2.8735632183908044%)  
 For year 1999: Fetching 200 samples out of 6056 (~3.3025099075297226%)  
 For year 2000: Fetching 200 samples out of 7466 (~2.678810608090008%)  
 For year 2001: Fetching 200 samples out of 8210 (~2.43605359317905%)  
 For year 2002: Fetching 200 samples out of 8061 (~2.4810817516437167%)  
 For year 2003: Fetching 200 samples out of 5872 (~3.4059945504087192%)  
 For year 2005: Fetching 200 samples out of 8269 (~2.4186721489902046%)  
 For year 2006: Fetching 200 samples out of 8810 (~2.2701475595913734%)  
 For year 2007: Fetching 200 samples out of 7863 (~2.543558438255119%)

For year 2009: Fetching 200 samples out of 17262 (~1.1586142973004288%)  
 For year 2010: Fetching 200 samples out of 11089 (~1.8035891423933628%)  
 For year 2011: Fetching 200 samples out of 13957 (~1.4329727018700293%)  
 For year 2012: Fetching 200 samples out of 16356 (~1.2227928588897041%)  
 For year 2013: Fetching 200 samples out of 10240 (~1.953125%)  
 For year 2014: Fetching 200 samples out of 12404 (~1.6123831022250887%)  
 For year 2016: Fetching 200 samples out of 8165 (~2.449479485609308%)  
 For year 2017: Fetching 200 samples out of 7270 (~2.751031636863824%)

Let's randomly choose a fixed number of documents (here currently: 200) from various different decades. Then passing (document, decade) pairs to the model below. The decade is computed by subtracting `mod(<year>, 10)` from `<year>`.

**Test data** Choose the other **2 years that were not selected** within the decades in the train set. In this case, **we do not have to limit the number of documents for year**. It doesn't make sense to undersample the test set since it represents "unseen" data. And, unseen data should be as close to reality as possible. That means, that it is normal that there are much more documents from later decades than from earlier.

So, instead of 200, we will put there a very large number to be sure that all the documents from every year are selected.

```
[30]: #seed for reproducibility
      random.seed(123)

      file_contents_test = []
      targets_test = []

      for year in [1913, 1917, 1930, 1939, 1950, 1954, 1976, 1977, 1994, 2013, 2015,
                  1923, 1927, 1940, 1949, 1960, 1964, 1986, 1987, 2004]:
          for file in get_files_for_year(year):
              file_contents_test.append(extract_words(file)[1])
              targets_test.append(year - year%10)
```

For year 1913: Fetching 2037 samples out of 2037 (~100.0%)  
 For year 1917: Fetching 0 samples out of 0 (~0%)  
 For year 1930: Fetching 30 samples out of 30 (~100.0%)  
 For year 1939: Fetching 1632 samples out of 1632 (~100.0%)  
 For year 1950: Fetching 3017 samples out of 3017 (~100.0%)  
 For year 1954: Fetching 1546 samples out of 1546 (~100.0%)  
 For year 1976: Fetching 2555 samples out of 2555 (~100.0%)  
 For year 1977: Fetching 2105 samples out of 2105 (~100.0%)  
 For year 1994: Fetching 8187 samples out of 8187 (~100.0%)  
 For year 2013: Fetching 10240 samples out of 10240 (~100.0%)  
 For year 2015: Fetching 18052 samples out of 18052 (~100.0%)  
 For year 1923: Fetching 26 samples out of 26 (~100.0%)  
 For year 1927: Fetching 59 samples out of 59 (~100.0%)  
 For year 1940: Fetching 1458 samples out of 1458 (~100.0%)  
 For year 1949: Fetching 2661 samples out of 2661 (~100.0%)

For year 1960: Fetching 1862 samples out of 1862 (~100.0%)  
For year 1964: Fetching 1677 samples out of 1677 (~100.0%)  
For year 1986: Fetching 3432 samples out of 3432 (~100.0%)  
For year 1987: Fetching 3212 samples out of 3212 (~100.0%)  
For year 2004: Fetching 9466 samples out of 9466 (~100.0%)

See classes distribution within train and test sets

```
[31]: from collections import Counter

print(Counter(targets).keys())
print(Counter(targets).values())

print(Counter(targets_test).keys())
print(Counter(targets_test).values())
```

```
dict_keys([1910, 1920, 1930, 1940, 1950, 1960, 1970, 1980, 1990, 2000, 2010])
dict_values([539, 218, 574, 1600, 1600, 1600, 1600, 1600, 1600, 1600, 1400])
dict_keys([1910, 1930, 1950, 1970, 1990, 2010, 1920, 1940, 1960, 1980, 2000])
dict_values([2037, 1662, 4563, 4660, 8187, 28292, 85, 4119, 3539, 6644, 9466])
```

We see that although there are some differences within the classes for the train split, it is acceptable to perform the classification task. Maximum within the classes for training is 1600.

Test set is expected to have much more class imbalance. However, our model should deal with it thanks to the undersampling that was performed.

### 3.5.2 Text feature extraction

We have considered 3 different methods for text feature extraction: Tf-idf, word2vec and doc2vec. All of them will be implemented through the corresponding functions from *sklearn* library.

**TF-IDF** Helper function to transform the data so that it is in the right format for the `tfidfVectorizer()` function that will be used later on:

```
[32]: class JoinElement(object):
        def fit(self, X, y):
            return self

        def transform(self, X):
            #joins the elements of a list (which represents a document) into a
            →single string
            #with a blank space separation between each word
            return [' '.join(X[i]) for i in range(len(X))]
```

More information about it: [sklearn documentation](#).

**Word2Vec** Original paper

: Mikolov, T., Sutskever, I., Chen, K., Corrado, G. S., & Dean, J. (2013). Distributed representations of words and phrases and their compositionality. *Advances in neural information processing systems*, 26, 3111-3119.

With this model every word is assigned a unique vector of configurable cardinality such that the dot product of two randomly chosen vectors should be proportional to the semantic similarity for the associated words. This happens during the training step using logistic regression and sliding windows. Personally I found that this video delivers a solid explanation of the concepts: <https://www.youtube.com/watch?v=QyrUentbkvw>

However, since we are working with entire documents as training items we have to somehow aggregate the vectors for every word in a given document. This can be done e.g. by taking the mean and/or summing up the vectors (see `MeanEmbeddingVectorizer`), optionally weighted by TF-IDF (see `MeanEmbeddingVectorizerTfidf`).

```
[33]: class MeanEmbeddingVectorizer(BaseEstimator):
    def fit(self, X, y):
        self.word2vec = Word2Vec(X)
        return self

    def transform(self, X):
        return np.array([
            np.mean([self.word2vec.wv[w] for w in words if w in self.word2vec.
↪wv.vocab]
                        or [np.zeros(self.word2vec.vector_size)], axis=0)
            for words in X
        ])

    def fit_transform(self, X, y):
        self.fit(X, y)
        return self.transform(X)
```

```
[34]: class MeanEmbeddingVectorizerTfidf(BaseEstimator):
    def fit(self, X, y):
        self.word2vec = Word2Vec(X)
        self.X_joined = [' '.join(X[i]) for i in range(len(X))]
        self.vectorizer = TfidfVectorizer()
        self.transformed = self.vectorizer.fit_transform(self.X_joined)
        self.transformed = pd.DataFrame.sparse.from_spmatrix(self.transformed)
        return self

    def tfidf(self, w, docid):
        if w in self.vectorizer.vocabulary_:
            return self.transformed[self.vectorizer.vocabulary_[w]][docid]
        else:
            return 0

    def transform(self, X):
```

```

        return np.array([
            np.mean([self.word2vec.wv[w] * self.tfidf(w, i) for w in words if w_
→in self.word2vec.wv.vocab]
                    or [np.zeros(self.word2vec.vector_size)], axis=0)
            for i, words in enumerate(X)
        ])

    def fit_transform(self, X, y):
        self = self.fit(X, y)
        return self.transform(X)

```

**Doc2Vec** Finally we are attempting to build a model using *Doc2Vec*. After training this model with our training corpus we receive a vector of configurable cardinality for each document.

Original paper

: Le, Quoc, and Tomas Mikolov. “Distributed representations of sentences and documents.” International conference on machine learning. 2014.

```

[35]: class Doc2Vectorizer(BaseEstimator):
    def __init__(self, window=2, vector_size=100):
        self.window = window
        self.vector_size = vector_size

    def fit(self, X, y):
        docs = [TaggedDocument(X[i], [y[i]]) for i in range(len(X))]
        self.doc_vec = Doc2Vec(docs, vector_size=self.vector_size, window=self.
→window, min_count=1, workers=4)
        return self

    def transform(self, X):
        return [self.doc_vec.infer_vector(X[i]) for i in range(len(X))]

```

**BERT** (*Bidirectional Encoder Representations from Transformers*) is also interesting to look at, but we’ll skip this here because we predict training a model from scratch would use up too many resources. Given more time however you could search for pretrained networks that roughly serve the purpose of classification of documents according to publication year.

Paper

: Devlin, J., Chang, M. W., Lee, K., & Toutanova, K. (2018). Bert: Pre-training of deep bidirectional transformers for language understanding. arXiv preprint arXiv:1810.04805.

### 3.5.3 Classifiers

3 different classifiers are going to be trained: Multinomial Naive Bayes, Support Vector Machines and Random Forest Classifier. All of them will be implemented using sklearn library.

**Multinomial Naive Bayes (MNB)** MNB is a common method for document classification due to its good balance between computational efficiency and predictive performance (Eibe, 2006). Therefore, we decided to choose it as one of our classifiers.

Details on the algorithm implementation can be found in the [sklearn documentation](#)..

The set of hyperparameters chosen to perform the grid search cross-validation during the training are based on the recommendations [from this article](#).

**Support Vector Machines** Support vector machines are widely used for classification purposes. What is more, it improves Multinomial Naive Bayes in terms of performance in most of the classification tasks. Thus, it was also chosen as one of our classifiers to be trained.

Details on the algorithm implementation can be found in [sklearn documentation](#).

The set of hyperparameters chosen to perform the grid search cross-validation during the training are based on the recommendations [from this article](#).

**Random Forest Classifier** Random Forest Classifier is one of the best methods according to the literature for classification tasks. However, the runtime may be extremely large (specially when increasing the size of the forest within grid search CV setups).

Details on the algorithm implementation can be found in [sklearn documentation](#)..

The set of hyperparameters chosen to perform the grid search cross-validation during the training are based on the recommendations [from this article](#)..

### 3.5.4 Train models

Since there are 3 methods for feature extraction and 3 classifiers, we should train 9 kind of models with their different combinations of hyperparameters. However, multinomial naive bayes does not take negative values produced by Word2Vec and Doc2Vec. Therefore, we have 7.

For each model, **a grid search is performed with different combinations of hyperparameters** for the classifiers and the text extraction methods. Afterwards, the most relevant results of each of the models are stored in a pandas data frame.

The goal of this grid search is to find the best combination of hyperparameters for each of our 7 combinations.

Note that **ideally we should perform a random search prior to the grid search to limit the scope of the best hyperparameters** to be used and then perform a more accurate search. However, this would lead to a tedious notebook and extremely large runtimes.

#### Model 1: TFIDF vectorizer, select K best and Multinomial Naive Bayes

```
[36]: #choose parameters for the different steps in the pipeline
parameters_model_1 = {

    #select KBest
    #k number of top features to select (default 10)
    "k_best_k": [10, 500],
    #score function to be used (default f_classif)
```



```

"k_best__score_func": [chi2],

#MultinomialNaiveBayes
#alpha is a parameter for smoothing (default value is 1)
"MNB__alpha": np.linspace(0.5, 1.5, 4),
#whether to learn class prior probabilities or not (default value is True)
"MNB__fit_prior": [True, False],

#TFIDF Vectorizer
#Whether the feature should be made of word or character n-grams (default_
→word)
"tfidf__analyzer": ["word"],
#Smooth idf weights by adding one to document frequencies,
#as if an extra document was seen containing every term in the collection_
→exactly once.
#Prevents zero divisions (default True)
"tfidf__smooth_idf": [True, False]

}

#build a pipeline
model_1_pipeline = Pipeline([
    #joins list into a single string
    ('join', JoinElement()),
    #tfidf vectorizer
    ('tfidf', TfidfVectorizer()),
    #select 1000 best from word vectors
    ('k_best', SelectKBest()),
    #apply naive bayes
    ('MNB', MultinomialNB())
])

#design grid search
grid_search_model_1 = GridSearchCV(
    #pipeline to be followed
    model_1_pipeline,
    #parameters
    param_grid=parameters_model_1,
    #number of folds for CV
    cv=5,
    #scoring to be considered for the cv
    scoring = "accuracy",
    #parallelize if possible
    n_jobs=-1
)

```

```

#fit the grid search for training data
grid_search_model_1.fit(file_contents, targets)

#save results of cross validation
cv_results_model_1 = pd.DataFrame(grid_search_model_1.cv_results_)

#filter columns to be kept in the dataframe
filter_col = [col for col in cv_results_model_1 if (col.startswith("param_") or
↳col.startswith("mean_") or col.startswith("rank"))]

#save results with only filtered columns
cv_results_model_1 = cv_results_model_1[filter_col]

#save name of the model for later comparison
cv_results_model_1.insert(loc=0, column="Model", value= "1")

#round mean_test_score
cv_results_model_1["mean_test_score"] = cv_results_model_1["mean_test_score"].
↳round(2)

#show best 5 sorted by mean_test_score
display(cv_results_model_1.sort_values(by="mean_test_score", ascending=False).
↳head(5))

```

	Model	mean_fit_time	mean_score_time	param_MNB__alpha	\
6	1	8.387684	2.188630	0.5	
7	1	8.919221	1.888311	0.5	
15	1	8.542751	1.953357	0.833333	
14	1	7.195633	1.739061	0.833333	
23	1	7.141477	2.184534	1.166667	

	param_MNB__fit_prior	param_k_best_k	param_k_best_score_func	\
6	False	500	<function chi2 at 0x7fac4119dc10>	
7	False	500	<function chi2 at 0x7fac4119dc10>	
15	False	500	<function chi2 at 0x7fac4119dc10>	
14	False	500	<function chi2 at 0x7fac4119dc10>	
23	False	500	<function chi2 at 0x7fac4119dc10>	

	param_tfidf__analyzer	param_tfidf__smooth_idf	mean_test_score	\
6	word	True	0.49	
7	word	False	0.49	
15	word	False	0.48	
14	word	True	0.48	
23	word	False	0.46	

	rank_test_score
--	-----------------

6	1
7	2
15	4
14	3
23	8

## Model 2: TFIDF vectorizer, select K best and SVC

```
[37]: #choose parameters for the different steps in the pipeline
parameters_model_2 = {

    #select KBest
    #k number of top features to select (default 10)
    "k_best_k": [10, 500],
    #score function to be used (default f_classif)
    "k_best_score_func": [chi2],

    #SVC
    #Specifies the kernel type to be used in the algorithm
    "SVC_kernel" : ["linear", "poly", "sigmoid"],
    #Kernel coefficient for 'rbf', 'poly' and 'sigmoid'
    "SVC_gamma": [1,0.1,0.001],
    #Regularization parameter. The strength of the regularization is inversely
    ↪proportional to C.
    #Must be strictly positive. The penalty is a squared l2 penalty.
    "SVC_C": [0.1,1, 10, 100],

    #TFIDF Vectorizer
    #Whether the feature should be made of word or character n-grams (default
    ↪word)
    "tfidf_analyzer": ["word"],
    #Smooth idf weights by adding one to document frequencies,
    #as if an extra document was seen containing every term in the collection
    ↪exactly once.
    #Prevents zero divisions (default True)
    "tfidf_smooth_idf": [True, False]

}

#build a pipeline
model_2_pipeline = Pipeline([
    #joins list into a single string
    ('join', JoinElement()),
    #tfidf vectorizer
    ('tfidf', TfidfVectorizer()),
    #select 1000 best from word vectors
    ('k_best', SelectKBest()),
```

```

        #apply naive bayes
        ('SVC', SVC())
    ])

#design grid search
grid_search_model_2 = GridSearchCV(
    #pipeline to be followed
    model_2_pipeline,
    #parameters
    param_grid=parameters_model_2,
    #number of folds for CV
    cv=5,
    #scoring to be considered for the cv
    scoring = "accuracy",
    #parallelize if possible
    n_jobs=-1
)

#fit the grid search for training data
grid_search_model_2.fit(file_contents, targets)

#save results of cross validation
cv_results_model_2 = pd.DataFrame(grid_search_model_2.cv_results_)

#filter columns to be kept in the dataframe
filter_col = [col for col in cv_results_model_2 if (col.startswith("param_") or
    ↳col.startswith("mean_") or col.startswith("rank"))]

#save results with only filtered columns
cv_results_model_2 = cv_results_model_2[filter_col]

#save name of the model for later comparison
cv_results_model_2.insert(loc=0, column="Model", value= "2")

#round mean_test_score
cv_results_model_2["mean_test_score"] = cv_results_model_2["mean_test_score"].
    ↳round(2)

#show best 5 sorted by mean_test_score
display(cv_results_model_2.sort_values(by="mean_test_score", ascending=False).
    ↳head(5))

```

	Model	mean_fit_time	mean_score_time	param_SVC__C	param_SVC__gamma	\
111	2	45.537666	12.704505	100	1	
135	2	49.423878	14.831820	100	0.001	
123	2	46.233693	13.522951	100	0.1	
134	2	37.751620	16.068168	100	0.001	

110	2	36.764385	14.584715	100	1
	param_SVC__kernel	param_k_best__k		param_k_best__score_func	\
111	linear	500		<function chi2 at 0x7fac4119dc10>	
135	linear	500		<function chi2 at 0x7fac4119dc10>	
123	linear	500		<function chi2 at 0x7fac4119dc10>	
134	linear	500		<function chi2 at 0x7fac4119dc10>	
110	linear	500		<function chi2 at 0x7fac4119dc10>	
	param_tfidf__analyzer	param_tfidf__smooth_idf		mean_test_score	\
111	word	False		0.54	
135	word	False		0.54	
123	word	False		0.54	
134	word	True		0.53	
110	word	True		0.53	
	rank_test_score				
111	1				
135	1				
123	1				
134	4				
110	4				

### Model 3: TFIDF vectorizer, select K best and Random Forest Classifier

```
[38]: #choose parameters for the different steps in the pipeline
parameters_model_3 = {

    #select KBest
    #k number of top features to select (default 10)
    "k_best__k": [10, 500],
    #score function to be used (default f_classif)
    "k_best__score_func": [chi2],

    #RF classifier
    #The number of trees in the forest (default is 100)
    "clf__n_estimators": [10,100,200],
    #The minimum number of samples required to split an internal node (default
    → is 2 but it is a large dataset)
    "clf__min_samples_split": [10, 40, 80],
    #The number of features to consider when looking for the best split
    → (default "auto" but sparse dataset)
    "clf__max_features": ["auto", 10],
    #maximum depth of the tree (default None)
    "clf__max_depth": [10, None],
```

```

#TFIDF Vectorizer
#Whether the feature should be made of word or character n-grams (default
→word)
"tfidf__analyzer": ["word"],
#Smooth idf weights by adding one to document frequencies,
#as if an extra document was seen containing every term in the collection
→exactly once.
#Prevents zero divisions (default True)
"tfidf__smooth_idf": [True, False]

}

#build a pipeline
model_3_pipeline = Pipeline([
    #joins list into a single string
    ('join', JoinElement()),
    #tfidf vectorizer
    ('tfidf', TfidfVectorizer()),
    #select 1000 best from word vectors
    ('k_best', SelectKBest()),
    #apply naive bayes
    ('clf', RandomForestClassifier())
])

#design grid search
grid_search_model_3 = GridSearchCV(
    #pipeline to be followed
    model_3_pipeline,
    #parameters
    param_grid=parameters_model_3,
    #number of folds for CV
    cv=5,
    #scoring to be considered for the cv
    scoring = "accuracy",
    #parallelize if possible
    n_jobs=-1
)

#fit the grid search for training data
grid_search_model_3.fit(file_contents, targets)

#save results of cross validation
cv_results_model_3 = pd.DataFrame(grid_search_model_3.cv_results_)

```

```

#filter columns to be kept in the dataframe
filter_col = [col for col in cv_results_model_3 if (col.startswith("param_") or
↳col.startswith("mean_") or col.startswith("rank"))]

#save results with only filtered columns
cv_results_model_3 = cv_results_model_3[filter_col]

#save name of the model for later comparison
cv_results_model_3.insert(loc=0, column="Model", value= "3")

#round mean_test_score
cv_results_model_3["mean_test_score"] = cv_results_model_3["mean_test_score"].
↳round(2)

#show best 5 sorted by mean_test_score
display(cv_results_model_3.sort_values(by="mean_test_score", ascending=False).
↳head(5))

```

	Model	mean_fit_time	mean_score_time	param_clf__max_depth	\
119	3	36.584321	3.331602	None	
143	3	17.497095	2.516746	None	
142	3	19.232731	2.954226	None	
131	3	27.325266	3.639129	None	
130	3	23.597651	3.784672	None	

	param_clf__max_features	param_clf__min_samples_split	\
119	10	10	
143	10	80	
142	10	80	
131	10	40	
130	10	40	

	param_clf__n_estimators	param_k_best__k	\
119	200	500	
143	200	500	
142	200	500	
131	200	500	
130	200	500	

	param_k_best__score_func	param_tfidf__analyzer	\
119	<function chi2 at 0x7fac4119dc10>	word	
143	<function chi2 at 0x7fac4119dc10>	word	
142	<function chi2 at 0x7fac4119dc10>	word	
131	<function chi2 at 0x7fac4119dc10>	word	
130	<function chi2 at 0x7fac4119dc10>	word	

	param_tfidf__smooth_idf	mean_test_score	rank_test_score
--	-------------------------	-----------------	-----------------

119	False	0.51	1
143	False	0.50	2
142	True	0.50	8
131	False	0.50	3
130	True	0.50	5

#### Model 4: Word2Vec and SVC

```
[39]: #choose parameters for the different steps in the pipeline
parameters_model_4 = {

    #SVC
    #Specifies the kernel type to be used in the algorithm
    "SVC__kernel" : ["linear", "poly", "sigmoid"],
    #Kernel coefficient for 'rbf', 'poly' and 'sigmoid'
    "SVC__gamma": [1,0.1,0.001],
    #Regularization parameter. The strength of the regularization is inversely
    →proportional to C.
    #Must be strictly positive. The penalty is a squared l2 penalty.
    "SVC__C": [0.1,1, 10, 100]

    #defaults for Word2Vec

}

#build a pipeline
model_4_pipeline = Pipeline([
    #tfidf vectorizer
    ('word2vec', MeanEmbeddingVectorizer()),
    #apply naive bayes
    ('SVC', SVC())
])

#design grid search
grid_search_model_4 = GridSearchCV(
    #pipeline to be followed
    model_4_pipeline,
    #parameters
    param_grid=parameters_model_4,
    #number of folds for CV
    cv=5,
    #scoring to be considered for the cv
    scoring = "accuracy",
    #parallelize if possible
    n_jobs=-1
)
```



```

#fit the grid search for training data
grid_search_model_4.fit(file_contents, targets)

#save results of cross validation
cv_results_model_4 = pd.DataFrame(grid_search_model_4.cv_results_)

#filter columns to be kept in the dataframe
filter_col = [col for col in cv_results_model_4 if (col.startswith("param_") or
↳col.startswith("mean_") or col.startswith("rank"))]

#save results with only filtered columns
cv_results_model_4 = cv_results_model_4[filter_col]

#save name of the model for later comparison
cv_results_model_4.insert(loc=0, column="Model", value= "4")

#round mean_test_score
cv_results_model_4["mean_test_score"] = cv_results_model_4["mean_test_score"].
↳round(2)

#show best 5 sorted by mean_test_score
display(cv_results_model_4.sort_values(by="mean_test_score", ascending=False).
↳head(5))

```

	Model	mean_fit_time	mean_score_time	param_SVC__C	param_SVC__gamma	\
31	4	343.308681	34.224248	100	0.1	
22	4	288.424709	26.118967	10	0.1	
18	4	278.652019	34.882533	10	1	
1	4	389.094729	50.654311	0.1	1	
27	4	497.729028	27.616480	100	1	

	param_SVC__kernel	mean_test_score	rank_test_score
31	poly	0.42	2
22	poly	0.42	1
18	linear	0.41	4
1	poly	0.41	3
27	linear	0.41	8

## Model 5: Word2Vec and Random Forest Classifier

```

[40]: #choose parameters for the different steps in the pipeline
parameters_model_5 = {

    #RF classifier
    #The number of trees in the forest (default is 100)
    "clf__n_estimators" : [10,100,200],
    #The minimum number of samples required to split an internal node (default_
↳is 2 but it is a large dataset)

```

```

    "clf__min_samples_split": [10, 40, 80],
    #The number of features to consider when looking for the best split
    → (default "auto" but sparse dataset)
    'clf__max_features': ["auto", 10],
    #maximum depth of the tree (default None)
    'clf__max_depth': [10, None]

    #defaults word2vec

}

#build a pipeline
model_5_pipeline = Pipeline([
    #tfidf vectorizer
    ('word2vec', MeanEmbeddingVectorizer()),
    #apply naive bayes
    ('clf', RandomForestClassifier())
])

#design grid search
grid_search_model_5 = GridSearchCV(
    #pipeline to be followed
    model_5_pipeline,
    #parameters
    param_grid=parameters_model_5,
    #number of folds for CV
    cv=5,
    #scoring to be considered for the cv
    scoring = "accuracy",
    #parallelize if possible
    n_jobs=-1
)

#fit the grid search for training data
grid_search_model_5.fit(file_contents, targets)

#save results of cross validation
cv_results_model_5 = pd.DataFrame(grid_search_model_5.cv_results_)

#filter columns to be kept in the dataframe
filter_col = [col for col in cv_results_model_5 if (col.startswith("param_") or
    → col.startswith("mean_") or col.startswith("rank"))]

#save results with only filtered columns
cv_results_model_5 = cv_results_model_5[filter_col]

```

```

#save name of the model for later comparison
cv_results_model_5.insert(loc=0, column="Model", value= "5")

#round mean_test_score
cv_results_model_5["mean_test_score"] = cv_results_model_5["mean_test_score"].
↳round(2)

#show best 5 sorted by mean_test_score
display(cv_results_model_5.sort_values(by="mean_test_score", ascending=False).
↳head(5))

```

	Model	mean_fit_time	mean_score_time	param_clf__max_depth	\
20	5	177.896227	9.965774	None	
32	5	154.525328	8.964136	None	
29	5	163.433541	8.489580	None	
35	5	113.754320	2.856136	None	
13	5	155.250164	10.702845		10

	param_clf__max_features	param_clf__min_samples_split	\
20	auto		10
32	10		40
29	10		10
35	10		80
13	10		40

	param_clf__n_estimators	mean_test_score	rank_test_score
20	200	0.33	1
32	200	0.32	2
29	200	0.32	3
35	200	0.31	15
13	100	0.31	18

## Model 6: Doc2Vec and Support Vector Machines

```

[41]: #choose parameters for the different steps in the pipeline
parameters_model_6 = {

    #SVC
    #Specifies the kernel type to be used in the algorithm
    "SVC__kernel" : ["linear", "poly", "sigmoid"],
    #Kernel coefficient for 'rbf', 'poly' and 'sigmoid'
    "SVC__gamma": [1,0.1,0.001],
    #Regularization parameter. The strength of the regularization is inversely
    ↳proportional to C.
    #Must be strictly positive. The penalty is a squared l2 penalty.
    "SVC__C": [0.1,1, 10, 100]

    #defaults doc2vec

```

```

}

#build a pipeline
model_6_pipeline = Pipeline([
    #tfidf vectorizer
    ('doc2vec', Doc2Vectorizer()),
    #apply naive bayes
    ('SVC', SVC())
])

#design grid search
grid_search_model_6 = GridSearchCV(
    #pipeline to be followed
    model_6_pipeline,
    #parameters
    param_grid=parameters_model_6,
    #number of folds for CV
    cv=5,
    #scoring to be considered for the cv
    scoring = "accuracy",
    #parallelize if possible
    n_jobs=-1
)

#fit the grid search for training data
grid_search_model_6.fit(file_contents, targets)

#save results of cross validation
cv_results_model_6 = pd.DataFrame(grid_search_model_6.cv_results_)

#filter columns to be kept in the dataframe
filter_col = [col for col in cv_results_model_6 if (col.startswith("param_") or
    ↳col.startswith("mean_") or col.startswith("rank"))]

#save results with only filtered columns
cv_results_model_6 = cv_results_model_6[filter_col]

#save name of the model for later comparison
cv_results_model_6.insert(loc=0, column="Model", value= "6")

#round mean_test_score
cv_results_model_6["mean_test_score"] = cv_results_model_6["mean_test_score"].
    ↳round(2)

#show best 5 sorted by mean_test_score

```

```
display(cv_results_model_6.sort_values(by="mean_test_score", ascending=False).
↳head(5))
```

	Model	mean_fit_time	mean_score_time	param_SVC__C	param_SVC__gamma	\
18	6	273.025912	28.329361	10	1	
12	6	245.849821	29.302286	1	0.1	
33	6	431.278508	23.150362	100	0.001	
3	6	228.663433	26.545853	0.1	0.1	
27	6	441.331615	27.363341	100	1	

	param_SVC__kernel	mean_test_score	rank_test_score
18	linear	0.43	4
12	linear	0.43	2
33	linear	0.43	8
3	linear	0.43	9
27	linear	0.43	6

### Model 7: Doc2Vec and Random Forest Classifier

```
[42]: #choose parameters for the different steps in the pipeline
parameters_model_7 = {

    #RF classifier
    #The number of trees in the forest (default is 100)
    "clf__n_estimators" : [10,100,200],
    #The minimum number of samples required to split an internal node (default
    ↳is 2 but it is a large dataset)
    "clf__min_samples_split": [10, 40, 80],
    #The number of features to consider when looking for the best split
    ↳(default "auto" but sparse dataset)
    'clf__max_features': ["auto", 10],
    #maximum depth of the tree (default None)
    'clf__max_depth': [10, None]

}

#build a pipeline
model_7_pipeline = Pipeline([
    #tfidf vectorizer
    ('doc2vec', Doc2Vectorizer()),
    #apply naive bayes
    ('clf', RandomForestClassifier())
])

#design grid search
grid_search_model_7 = GridSearchCV(
```

```

#pipeline to be followed
model_7_pipeline,
#parameters
param_grid=parameters_model_7,
#number of folds for CV
cv=5,
#scoring to be considered for the cv
scoring = "accuracy",
#parallelize if possible
n_jobs=-1
)

#fit the grid search for training data
grid_search_model_7.fit(file_contents, targets)

#save results of cross validation
cv_results_model_7 = pd.DataFrame(grid_search_model_7.cv_results_)

#filter columns to be kept in the dataframe
filter_col = [col for col in cv_results_model_7 if (col.startswith("param_") or
→col.startswith("mean_") or col.startswith("rank"))]

#save results with only filtered columns
cv_results_model_7 = cv_results_model_7[filter_col]

#save name of the model for later comparison
cv_results_model_7.insert(loc=0, column="Model", value= "7")

#round mean_test_score
cv_results_model_7["mean_test_score"] = cv_results_model_7["mean_test_score"].
→round(2)

#show best 5 sorted by mean_test_score
display(cv_results_model_7.sort_values(by="mean_test_score", ascending=False).
→head(5))

```

	Model	mean_fit_time	mean_score_time	param_clf__max_depth	\
23	7	252.317726	21.616184		None
20	7	260.065633	21.306356		None
28	7	255.810332	26.719237		None
29	7	269.841011	22.845155		None
19	7	242.370124	21.902374		None

	param_clf__max_features	param_clf__min_samples_split	\
23	auto	40	
20	auto	10	
28	10	10	

29	10	10
19	auto	10

	param_clf__n_estimators	mean_test_score	rank_test_score
23	200	0.37	1
20	200	0.37	2
28	100	0.36	5
29	200	0.36	4
19	100	0.36	7

### 3.5.5 Compare CV results from trained models

In this section, the results from CV are compared within the trained models.

**Raw results** Export results from grid search. This allows us to experiment with visualizations and results from CV without having to rerun the whole script.

```
[43]: #cv_results.to_csv("cv_results.csv",index=False)
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-43-d3b59e1ef88e> in <module>
----> 1 cv_results.to_csv("cv_results.csv",index=False)

NameError: name 'cv_results' is not defined
```

A dataframe showing the best models according to the **mean accuracy within the test folds** used for cross validation.

```
[45]: #merge cv results into 1 that keeps the relevant information

#empty dataframe that will keep all the results
cv_results = pd.DataFrame()

#loop over cv results
for i in [cv_results_model_1, cv_results_model_2, cv_results_model_3,
cv_results_model_4,
cv_results_model_5, cv_results_model_6, cv_results_model_7]:

    #select relevant columns
    selected = i[["Model", "mean_fit_time", "mean_score_time", "mean_test_score"]]

    #append to cv results
    cv_results = cv_results.append(selected)

#show models with best scores
display(cv_results.sort_values(by="mean_test_score", ascending=False).head(20))
```

```
cv_results.to_csv("cv_results2.csv", index=False)
```

	Model	mean_fit_time	mean_score_time	mean_test_score
111	2	45.537666	12.704505	0.54
135	2	49.423878	14.831820	0.54
123	2	46.233693	13.522951	0.54
122	2	37.660268	16.118198	0.53
110	2	36.764385	14.584715	0.53
134	2	37.751620	16.068168	0.53
87	2	50.362453	13.930855	0.52
74	2	42.097400	17.334257	0.52
75	2	53.180003	15.568920	0.52
86	2	39.190326	16.937734	0.52
98	2	22.771912	9.259802	0.52
99	2	28.150809	8.853756	0.52
131	2	55.788226	16.034667	0.52
130	2	42.225446	18.179445	0.52
114	2	42.626273	18.166295	0.51
119	3	36.584321	3.331602	0.51
115	2	55.720585	15.588012	0.51
118	3	33.057600	3.987268	0.50
131	3	27.325266	3.639129	0.50
115	3	23.073168	3.311102	0.50

*This table could be improved by also indicating the parameters used in each model but I thought it would be a bit overwhelming*


Models x y seem to achieve a better accuracy since they appear more often within the first positions.

**Tradeoff score vs mean fit time** A plot to check if there is some kind of tradeoff between accuracy and runtime of the algorithms.

```
[46]: #group by model
groups = cv_results.groupby("Model")

# Plot
fig, ax = plt.subplots()
ax.margins(0.05) # Optional, just adds 5% padding to the autoscaling
for name, group in groups:
    ax.plot(group.mean_fit_time, group.mean_test_score, marker='o',
            linestyle='', ms=12, label="Model %s" %name)
ax.legend(loc = 1)
plt.xlabel("Mean fit time")
plt.ylabel("Mean test Score (accuracy)")
plt.show()
```





Clear relation between runtime of the models and accuracy within test folds?


Some pre-processing or classifier takes more time to be run?

Further analysis on what it increases runtime of the models

### Best estimator from each model

```
[47]: #group by model and take best mean_test_score for each type of model
best_models = cv_results.groupby("Model")[["mean_test_score"]].max()

#plot
plt.bar(best_models.index, best_models["mean_test_score"])
plt.xlabel("Model")
plt.ylabel("Mean_test_score")
plt.show()
```



As expected best models are x and y

Worse models are x and y

Are they stable to changes in hyperparameters?

```
[48]: #mean of mean scores within the folds for each model
stability = cv_results.groupby("Model")[["mean_test_score"]].mean()
#standard deviation of the same
stability["standard_deviation"] = cv_results.
    ↳groupby("Model")[["mean_test_score"]].std()

#show
stability
```

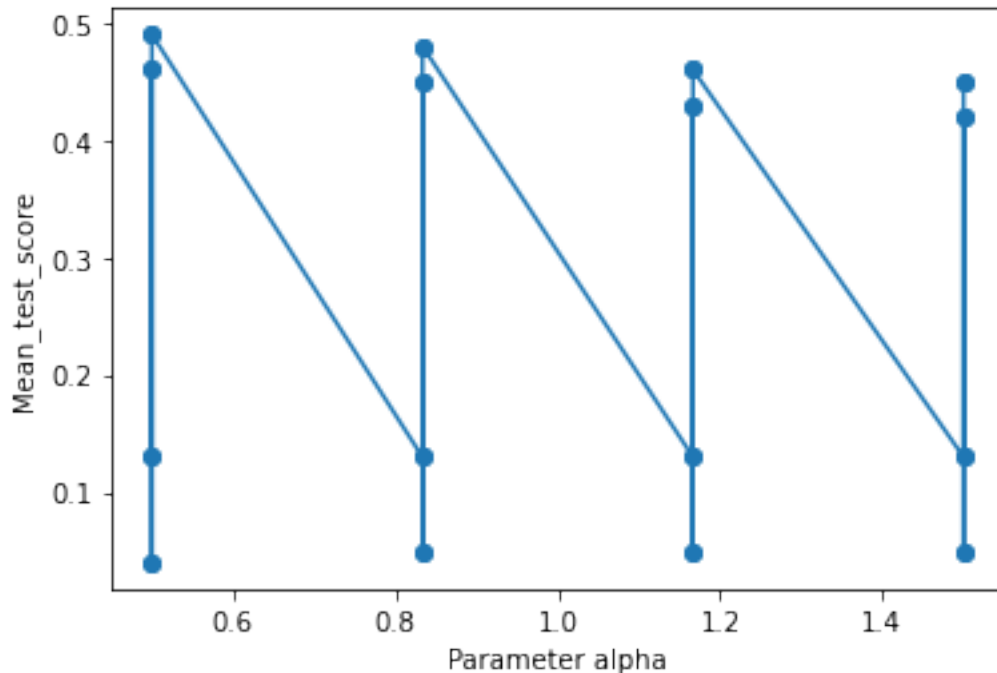
```
[48]:
```

Model	mean_test_score	standard_deviation
1	0.271875	0.189063
2	0.234861	0.139133
3	0.314028	0.143524
4	0.290556	0.128973
5	0.292778	0.024913
6	0.302222	0.140471
7	0.336944	0.023030

Model x is not very stable to the change of hyperparameters. What is influencing it so much?

Show an example

```
[49]: plt.plot(cv_results_model_1["param_MNB__alpha"],  
             ↪cv_results_model_1["mean_test_score"], marker="o")  
plt.xlabel("Parameter alpha")  
plt.ylabel("Mean_test_score")  
plt.show()
```



**Next steps** This grid searches helped us to see which combinations of models and hyperparameters are expected to be the best, how stable they are and other insights. However, **we cannot draw strong conclusions on this since we are still dealing with train data**. This step is only helping us to understand the models better and choose the ones with which we want to test (or validate). To proceed further we decided to select **the best combination of hyperparameters for each of the 7 models, predict on test data**, evaluate and draw conclusions. That is done in the next section of this notebook. We know that this are not strictly the best 7 models (see raw results), but we wanted to include more diversity.

### 3.6 Evaluation and model selection

**Predict on test data** using the best combinations of hyperparameters used in the models obtained in the training phase and evaluate using different metrics.

*This simulates predictions on unseen data. However, since it is done for many models and then we will choose the best model out of them, it behaves more like a validation set that would help us choose which model we would apply to actually unseen data.*

```

[50]: #add models to be evaluated
models = [
    grid_search_model_1,
    grid_search_model_2,
    grid_search_model_3,
    grid_search_model_4,
    grid_search_model_5,
    grid_search_model_6,
    grid_search_model_7
]

evaluation = pd.DataFrame(columns=["model"
                                , "mean_fit_time", "accuracy"
                                , "recall_macro", "recall_micro"
                                , "precision_macro", "precision_micro"
                                , "f1_macro", "f1_micro"
                                , "model_definition"
                                ])

i = -1 # Ensure that first item is index 0 in the loop
for model_ in models:
    # Yucky method of finding mean fit times:
    i = i + 1
    mean_fit_time = cv_results.groupby("Model")["mean_fit_time"].mean()[i]

    # Predict
    preds = model_.best_estimator_.predict(file_contents_test)
    model = cv_results.iloc[model_.best_index_,0]

    # Calculate metrics
    to_append = [
        "Model " + str(i+1),
        mean_fit_time,
        accuracy_score(y_true=targets_test,y_pred=preds),
        #choose micro or macro according to criteria
        recall_score(y_true=targets_test,y_pred=preds, average="macro"),
        recall_score(y_true=targets_test,y_pred=preds, average="micro"),
        precision_score(y_true=targets_test,y_pred=preds, average="macro"),
        precision_score(y_true=targets_test,y_pred=preds, average="micro"),
        f1_score(y_true=targets_test,y_pred=preds, average="macro"),
        f1_score(y_true=targets_test,y_pred=preds, average="micro"),
        model_
    ]

    #Append Metrics
    evaluation_length = len(evaluation)
    evaluation.loc[evaluation_length] = to_append

```

```

    # Print results and Confusion Matrix for each model
    print("#####")
    print("#####")
    print("
            Model "+ str(i+1) + ":")
    print("-----")

    print(evaluation.loc[i, 'mean_fit_time':'f1_macro'])
    print("\n")
    print("Pipeline: ")
    print(model_.best_estimator_)

    print("\n")
    print("Confusion Matrix: ")
    fig, ax = plt.subplots(figsize=(10, 10))
    plot_confusion_matrix(estimator=model_.best_estimator_,
                          , X=file_contents_test
                          , y_true=targets_test
                          , ax=ax
                          )

    plt.show()

# Print table of the models compared and sorted:
evaluation.sort_values(by="accuracy", ascending = False)

```

```

#####
#####
Model 1:

```

```

-----
mean_fit_time      8.04884
accuracy           0.580828
recall_macro       0.582249
recall_micro       0.580828
precision_macro     0.540098
precision_micro     0.580828
f1_macro           0.537704
Name: 0, dtype: object

```

Pipeline:


```

Pipeline(steps=[('join', <__main__.JoinElement object at 0x7fab1fd0be50>),
                 ('tfidf', TfidfVectorizer()),
                 ('k_best',
                  SelectKBest(k=500,

```

```
score_func=<function chi2 at 0x7fac4119dc10>)),
('MNB', MultinomialNB(alpha=0.5, fit_prior=False))])
```

Confusion Matrix:



```
#####
#####
```

Model 2:


```
-----
mean_fit_time    49.922866
accuracy         0.617782
recall_macro     0.605168
recall_micro     0.617782
precision_macro  0.588218
precision_micro  0.617782
```

```
f1_macro          0.579638
Name: 1, dtype: object
```

Pipeline:

```
Pipeline(steps=[('join', <__main__.JoinElement object at 0x7faaed06afa0>),
                 ('tfidf', TfidfVectorizer(smooth_idf=False)),
                 ('k_best',
                  SelectKBest(k=500,
                              score_func=<function chi2 at 0x7fac4119dc10>)),
                 ('SVC', SVC(C=100, gamma=1, kernel='linear'))])
```

Confusion Matrix:



#####

#####

Model 3:


-----  
mean\_fit\_time 13.328344  
accuracy 0.6174  
recall\_macro 0.594453  
recall\_micro 0.6174  
precision\_macro 0.621433  
precision\_micro 0.6174  
f1\_macro 0.588941  
Name: 2, dtype: object

Pipeline:

```
Pipeline(steps=[('join', <__main__.JoinElement object at 0x7fab1fd064f0>),  
                 ('tfidf', TfidfVectorizer(smooth_idf=False)),  
                 ('k_best',  
                  SelectKBest(k=500,  
                               score_func=<function chi2 at 0x7fac4119dc10>)),  
                 ('clf',  
                  RandomForestClassifier(max_features=10, min_samples_split=10,  
                                         n_estimators=200))])
```

Confusion Matrix:





#####  
#####

Model 4:

```
-----
mean_fit_time      320.122003
accuracy           0.494908
recall_macro       0.493105
recall_micro       0.494908
precision_macro    0.437746
precision_micro    0.494908
f1_macro           0.44268
Name: 3, dtype: object
```


Pipeline:

```

Pipeline(steps=[('word2vec', MeanEmbeddingVectorizer()),
                 ('SVC', SVC(C=10, gamma=0.1, kernel='poly'))])

```

Confusion Matrix:



```

#####
#####

```

Model 5:

```


-----
mean_fit_time      148.909811
accuracy           0.375611
recall_macro       0.33446
recall_micro       0.375611
precision_macro    0.391935
precision_micro    0.375611

```

```
f1_macro          0.330763
Name: 4, dtype: object
```

```
Pipeline:
Pipeline(steps=[('word2vec', MeanEmbeddingVectorizer()),
                 ('clf',
                  RandomForestClassifier(min_samples_split=10,
                                         n_estimators=200))])
```

Confusion Matrix:



```
#####
#####
Model 6:
```


---

```
mean_fit_time      276.770212
accuracy           0.465367
recall_macro       0.475932
recall_micro       0.465367
precision_macro     0.463425
precision_micro     0.465367
f1_macro           0.442343
Name: 5, dtype: object
```

Pipeline:

```
Pipeline(steps=[('doc2vec', Doc2Vecrizer()),
                 ('SVC', SVC(C=10, gamma=0.001, kernel='linear'))])
```

Confusion Matrix:



```
#####  
#####
```


Model 7:

```
-----  
mean_fit_time      241.609706  
accuracy           0.430543  
recall_macro       0.364123  
recall_micro       0.430543  
precision_macro     0.474474  
precision_micro     0.430543  
f1_macro           0.366073  
Name: 6, dtype: object
```

Pipeline:

```
Pipeline(steps=[('doc2vec', Doc2Vectorizer()),  
                 ('clf',  
                  RandomForestClassifier(min_samples_split=40,  
                                         n_estimators=200))])
```

Confusion Matrix:



```
[50]:      model  mean_fit_time  accuracy  recall_macro  recall_micro  \
1  Model 2      49.922866  0.617782    0.605168    0.617782
2  Model 3     13.328344  0.617400    0.594453    0.617400
0  Model 1      8.048840  0.580828    0.582249    0.580828
3  Model 4    320.122003  0.494908    0.493105    0.494908
5  Model 6    276.770212  0.465367    0.475932    0.465367
6  Model 7    241.609706  0.430543    0.364123    0.430543
4  Model 5    148.909811  0.375611    0.334460    0.375611

precision_macro  precision_micro  f1_macro  f1_micro  \
1      0.588218      0.617782  0.579638  0.617782
2      0.621433      0.617400  0.588941  0.617400
0      0.540098      0.580828  0.537704  0.580828
3      0.437746      0.494908  0.442680  0.494908
5      0.463425      0.465367  0.442343  0.465367
```

6	0.474474	0.430543	0.366073	0.430543
4	0.391935	0.375611	0.330763	0.375611

		model_definition
1	GridSearchCV(cv=5,\n	estimator=Pip...
2	GridSearchCV(cv=5,\n	estimator=Pip...
0	GridSearchCV(cv=5,\n	estimator=Pip...
3	GridSearchCV(cv=5,\n	estimator=Pip...
5	GridSearchCV(cv=5,\n	estimator=Pip...
6	GridSearchCV(cv=5,\n	estimator=Pip...
4	GridSearchCV(cv=5,\n	estimator=Pip...

The **best estimator we found after everything is:**

Main **conclusion:** for unseen data, we would choose to use the estimator from above.

## 4 Conclusion

**Draw conclusions from above. but add more evaluations and stuff.... Should probably be done after proper train/test is defined, and full model is tried**

Some observations: \* Model 1,2,3 (TFIDF) seems to work best overall, no matter training size  
 \* Model 5,6,7 (Doc2Vec) takes by far the most time, but their accuracy greatly increased when increasing training set

As often is the case in the fields of science, not all research leads to useable results. We ended up having to remodel our plans several times during this project, including a complete pivot of the datasets.

This did however give us some insight into how larger projects are managed. This also lead us to an interesting path of looking at a relatively obscure language.

Although further works is possible, we reached the conclusion that there is a change in the Icelandic spoken language throughout time, and it is therefore possible to train models that estimates which decade a given speech is from. However, take into account what explained in section 3.2: it may be also due to other factors (for instance, topic used).

Overall we did work with Data-Oriented Programming best practices. We were able to develop a scientific workflow. From the given data, we managed to train a model for prediction on test data with **decent results**.

## 5 Further Works

As we drilled down this dataset, we kept getting new ideas that we would like to experiment with, and try to gain better insight. Specifically, our next steps would be:

### 5.1 Predict Different Sources

As it currently stands, we are trying to estimate a decade of speeches from “Althingi”. However, the dataset has several other sources of Icelandic; both written and spoken (TV scripts, cinema

and others).

We would like to see if it was possible to extend our model to be able to classify the source.

## **5.2 Treating years as Continuous Variables**

We are currently treating decades as a class. By discretizing results from a regression algorithm, we think it should be possible to keep some nominal knowledge of the ordering of the years, and thus improving our predictions. It would be also interesting to see if we can achieve also decent prediction by narrowing a bit the intervals for the years (instead of decades, lustrums). And of course, it would be interesting to rerun the model in a more powerful machine using all the decades instead of discarding the intermediate ones.

## **5.3 Gaining insight into Explanatory Variables**

From our results, it is clear that it is somewhat possible to predict the decades. However, we are still treating the algorithms as “Black Boxes”. We would like to dive deeper into the decision trees/boundaries, to see if we can locate what it is that makes the predictions possible. It might be new words introduced, semantic changes, or something entirely different.

## **5.4 Additional Feature Extraction and Classifiers**

We would like to extend the list to include more classifiers, as well as trying to develop some additional feature extractions. E.g. “Glove Embedding” E.g. “Neural networks”